

3

• expressivism

As we saw at the end of the previous chapter, nonnaturalists generally assume that ethical language is *representational* and ethical thought is cognitive. This assumption is common to realist positions. The basic idea is that ethical statements and the states of mind that they express represent reality as being a particular way, and they are correct when reality is indeed this way. Expressivist theories in metaethics start out from skepticism about this idea. In order to argue that a representationalist understanding of ethical statements is mistaken, expressivists typically encourage us to begin our metaethical enquiry by asking not about (i) the nature of some kind of supposedly real thing: ethical value, but about (ii) the nature of one of our practices: making ethical evaluations. Then, their answer to question (ii) typically claims some interesting disanalogy between ethical evaluations and representations of reality.

In their philosophy of mind, this typically leads expressivists to endorse some sort of *non-cognitivism* about ethical thoughts. That is, they claim that the states of mind expressed by ethical statements have a desire-like rather than belief-like direction of fit with reality; they are something like desires, emotive reactions, preferences, or plans. By claiming that ethical sentences and the states of mind that they express are not in the business of representing reality, expressivists pave the way to an antirealist view, since they see no reason to commit to the reality of objective ethical facts. We call their view “expressivism” because the key claim that distinguishes it from the other main metaethical views is a claim about what kind of mental state is expressed by ethical statements.

In this chapter, we’ll discuss several increasingly sophisticated versions of this view and the objections that lead to increased sophistication. But we start with a few of the general reasons some philosophers have favored this view.

Key Point: Expressivism is a form of antirealism that understands ethical statements as expressions of evaluative attitudes rather than beliefs representing how reality is.

• ARGUMENTS IN FAVOR OF EXPRESSIVISM

The full case in favor of expressivism (or any other metaethical view) can only be pursued in comparison to the other theoretical options. So, in Chapter 7 we'll attempt to map many of the various dimensions of theoretical costs and benefits in metaethics. But before we get to that, it will prove helpful to canvass some of the initial motivations in favor of the view. Here, I start with three related motivations for expressivism.

Hume's Law and open-question intuitions + naturalistic worldview

As we saw before, Hume insisted on an important conceptual distinction between 'ought' and 'is,' and this has been pushed forward by various open-question style arguments in favor of the nonnaturalist view we encountered in the previous chapter that ethical facts are not reducible to natural facts. As Ayer (1936/1946, Chapter 6) pointed out, an equally valid response to Hume's Law and open-question intuitions is to argue that ethical thought and discourse aren't even in the business of representing aspects of reality. That is, rather than capturing the apparent difference in terms of the kinds of facts (natural vs. nonnatural) represented by ethical statements and the states of mind they express, one of the founding ideas of expressivism is that we can capture this in terms of the function of ethical statements and the states of mind they express. More specifically, assuming ordinary statements express beliefs about reality, expressivists claim that ethical statements express some noncognitive or non-representational state of mind—something like a desire, emotion, preference, or plan.

QUI: What distinguishes a representational state of mind from a nonrepresentational state of mind?

As far as that goes, we have two competing explanations of Hume's Law and open-question intuitions: the nonnaturalist's and the expressivist's. Which should we prefer? Expressivism is usually motivated by a naturalistic worldview. If you find Hume's Law and open-question intuitions persuasive, and you're also committed to a naturalistic worldview, then the expressivist's explanation is likely to look pretty good. It's not that expressivists *have to* deny that there are nonnatural facts (and that ethical statements are about such facts). But they will argue that once we have recognized the fundamentally different function of ethical statements and the states of mind they express, there will be no explanatory or theoretical need to posit an autonomous realm of nonnatural facts. The principle of parsimony encourages us to do without them.

Argument from motivational internalism + Humean theory of motivation

The intuition behind Hume's Law and open-question style arguments is not just that ethical statements and the states of mind they express are different, but also that they are different *because of how they bear on action*. The basic idea stems from the Humean belief–desire psychology of motivation. Recall from Chapter 2 that many philosophers have found it convenient to think of action as the product of two distinct kinds of mental states—states that represent the way reality is and states that set a goal for the agent to change reality. These are commonly (though somewhat misleadingly) called *beliefs* and *desires*.

Someone who thinks the Humean belief–desire psychology is basically right isn't thereby committed to any specific metaethical view. However, if one is also attracted to some form of motivational internalism about ethical thought, expressivism becomes increasingly tempting. Motivational internalism, recall, was the view that there is some especially tight connection between ethical thoughts and motivation. For example, thinking that one ought not steal seems to be tightly correlated with one's being motivated not to steal. Of course, sometimes we do things that we think we ought not do. But often these can be explained by our being overpowered by the countervailing motivations of temptation, anger, envy, etc.; or they can be explained by our not fully comprehending what we were doing when we acted wrongly. It would be strange if someone thought that she ought not ϕ but had absolutely no motivation to refrain from ϕ ing.

This provides the resources for a prominent argument in favor of expressivism. If the distinctive feature of ethical thoughts is that they are internally connected to motivation, and one thinks that beliefs alone cannot motivate, then a natural conclusion is that ethical thoughts aren't beliefs about the way reality is, but instead some desire-like state of mind. That is to say that expressivists sometimes argue that the reason that motivational internalism is true is that ethical thoughts can play the goal-setting role in a Humean explanation of motivation unlike ordinary empirical thoughts; and if this is right, then ethical statements are plausibly thought to express desire-like states of mind rather than beliefs about the way reality is.

What does it mean to call these states “desire-like”? It means that they are goal-setting rather than reality-representing. It may be reasonable to distinguish several different species of goal-setting mental states: ordinary desires (e.g., wanting to have some oysters), comparative preferences (e.g., preferring living in New York to Sydney), moral stances (e.g., abhorring racism), conditional plans (e.g., planning to save a drowning child if one sees one). An important further question for expressivists is which of these is their more specific account of the nature of moral judgments. But as far as the argument from motivational internalism is concerned, it just has to be some “desire-like” state of mind.

Key Point: Expressivism would provide a plausible explanation of how moral judgments could be internally connected to motivation provided that the Humean theory of motivation is correct.

Argument from supervenience

The final general motivation for expressivism has to do with supervenience. Recall that a liability of nonnaturalist views in metaethics is that they have no easy explanation of the way the ethical status of actions seems to depend intimately on the non-ethical facts of the situation. It would be very strange indeed if two actions that were identical in all relevant non-ethical respects were nonetheless evaluated differently ethically. For one action to be permissible and another action to be impermissible, for example, it seems that there must be some underlying non-ethical difference. For this reason many philosophers say that the ethical supervenes on the non-ethical.

This objection to nonnaturalism was originally pressed by expressivists such as Hare (1952) and Blackburn (1971), and they thought that expressivism offers a fairly compelling explanation of supervenience. The idea is to reject the nonnaturalist's picture of two distinct realms of fact (ethical and non-ethical), whose relation must be explained in such a way to secure supervenience. In its place, expressivists suggest that supervenience intuitions reflect a natural consistency constraint on our desire-like reactions to the way we take the reality to be. If, for instance, there are two circumstances that are exactly the same in all relevant descriptive ways, then approving of one while disapproving of the other would be arbitrary. And holding and expressing such arbitrary attitudes would undermine the natural coordinative function of having these attitudes. Hence, over time there would have been pressure towards respecting this consistency constraint in these attitudes. So much so, in fact, that the constraint would have become part of the meaning of the relevant ethical terms. According to the expressivist, this manifests itself in the intuitions of supervenience.

To get the idea, consider a nonmoral example of choosing pieces of cake. If one believes that two pieces of cake are alike in all relevant respects (same size, same kind, same taste, etc.), then it'd be hard to make sense of someone preferring one of the pieces to the other. To put it differently, if someone said, "I know they are exactly alike in respects x, y, z, but I want piece A instead of piece B," we'd be inclined to think that there was some other respect in which they are different that explains this difference in desires. The preferability of pieces of cake, it seems, supervenes on a range of descriptive features (size, kind, taste, etc.).

If morality supervenes on the non-ethical facts in this way, and ethical thoughts are desire-like as the expressivist contends, then we could offer the same sort of explanation of why any difference in ethical evaluation must, on pain of irrationality, rest on the assumption

that there is a relevant difference in non-ethical features. Again, we see an inference to the best explanation. The ethical seems to supervene on the non-ethical: expressivists argue that the best explanation of this is that ethical thoughts are desire-like reactions to how we take the non-ethical facts to be, not beliefs in some separate realm of nonnatural facts.

As far as that goes, however, the nonnaturalist might try to steal a leaf from the expressivist's playbook. She might argue that ethical beliefs are about nonnatural facts but then go on to insist that such beliefs are constrained by rational consistency in a similar way to the way expressivists think moral attitudes are constrained by rational consistency. That is to say that just like the expressivist holds that one's ethical view of two situations should not differ unless there is some underlying descriptive difference between the situations, perhaps the nonnaturalist could say something similar. In response, those sympathetic to expressivism will wonder why we should think that there are facts that are such that belief in them is constrained in this way. So this response on the part of the nonnaturalist would open them up to a metaphysical challenge, but they will think that this challenge can be met.

• VERSIONS AND OBJECTIONS

To appreciate the force of some of the objections that have been raised against expressivism, it will prove helpful to sketch a few more specific versions of the view. We'll use these sketches to explain the objections and then to explore the ways that expressivism has been refined to avoid the objections (though often with the consequence of inspiring new objections). Each subsequent view has been viewed as more plausible than the previous, in part due to reacting to objections, though some objections may remain.

Emotivism

Ayer's (1936/1946, Chapter 6) "Emotive Theory of Values" and Stevenson's "Emotive Meaning of Ethical Terms" (1937) are two key forbearers of more contemporary expressivist theories. Stevenson argued from ordinary language considerations that ethical words are primarily used to vent feelings, create moods, and incite people to action—rather than to record, clarify, and communicate beliefs. However, because it is not entirely clear whether to classify Stevenson as a proto-expressivist or a proto-relativist, I'll focus here on Ayer. Ayer's basic idea was that ethical statements should not be viewed as attempts to state matters of fact but rather to express emotive reactions to the facts (or at least what we take to be the facts).

To be more precise, we should say that Ayer argued that ethical words have at least two uses. First, they are used to describe the moralities of different communities of people; second, they are used to express one's own opinions about what is right, wrong, etc. According to Ayer, statements with ethical words used in the first way

are sociological descriptions rather than expressions of normative judgments. For instance, we might say something like, “In Saudi Arabia, it’s wrong for Muslim women to go bareheaded outside of their home,” not meaning to *condemn* Muslim women going bareheaded while in Saudi Arabia but to *describe* the moral sensibility of Saudi Arabian society. Ayer thought statements with ethical words used in the second way, however, are not descriptions of a moral sensibility, nor even descriptions of our own moral sensibilities. Rather they are direct expressions of the emotions comprising these sensibilities. For example, if one uses an ethical word in the second way saying, “Stealing is wrong,” Ayer’s view was that this expresses one’s negative attitude towards stealing rather than describing stealing or one’s attitudes towards stealing.

QU2: What’s the difference between describing yourself as feeling some way and expressing that feeling?

One helpful way to conceptualize this theory is as the “boo/hooray! theory” of ethical language. Ayer thought that saying that something is wrong or bad is not describing it as having a particular property, but rather expressing a negative emotion towards it—like booing it. And, similarly, saying that something is right or good is not describing it as having a particular property but rather expressing a positive emotion towards it—like hooraying it.

To understand why many philosophers have been dissatisfied with emotivism, it’s useful to delve into a bit of the history around the development of Ayer’s metaethical view. He was largely concerned with issues in the philosophy of language, where he defended the **verificationist principle** of meaning. This says that the factual meaning of a sentence is the conditions under which it would be “verified” or confirmed as true. For example, the verification conditions of a sentence such as, “Crocodiles live in Florida” would be the conditions under which one could see or infer the existence of crocodiles living in Florida (presumably there is some range of experiences that would confirm this as true).

Ayer used the verificationist principle of meaning as a kind of philosophical weapon, arguing that various philosophers’ views in metaphysics were unverifiable and so, strictly speaking, meaningless (since there would have to be some conditions under which they would be verified in order to have meaning, according to the verificationist principle of meaning). However, Ayer ran up against the fact that it often seems impossible to verify ethical claims. For instance, if someone says, “Manning’s leak of the classified documents was morally wrong,” how would we go about verifying this as true, or verifying its denial as false? It’s not like there is some range of experiences which would settle the matter.

This may have led some philosophers to go back and reconsider the verificationist theory of meaning, but Ayer stood fast to this principle and embraced metaethical conclusions that seemed to him to follow from it. I’ll mention four of these conclusions here.

First, Ayer argued that moral disagreement is senseless, unless it proceeds on the assumption of shared values. His idea was that there can be cases where two people share an ethical outlook, e.g., perhaps they both think economic conditions that lead to some people living in abject poverty are morally bad. Given such shared ethical values, there could be genuine disagreement about something like whether a particular governmental policy is morally bad. According to Ayer, however, this would actually be a factual disagreement and not a moral disagreement. True, we frame the disagreement as one about whether something is morally bad, but what the people disagree about is whether the governmental policy leads to abject poverty, not whether abject poverty is bad. But if the two parties to the disagreement didn't share ethical values, then it would be silly for them to disagree. In that case, according to Ayer, there's nothing really for them to disagree about—they simply have different ethical viewpoints. Accordingly, he thought that there is no such thing as genuinely *ethical* disagreement. Apparent ethical disagreements are either factual disagreements on the backdrop of shared values or simply people not sharing ethical values.

Second, Ayer argued that there's actually not much for moral philosophers to do, as there's not really a subject matter for them to figure out (remember he thought that there is no way to verify many ethical claims as true or false). He thought that different people have different values (which is something any anthropologist or sociologist could tell us), and all there is for the moral philosopher to do is to note that ethical words have no factual meaning. The remaining task of describing the different feelings ethical words are used to express, Ayer thought, is a matter for psychology not philosophy.

Third (and most famously), Ayer claimed that ethical statements "...do not come under the category of truth and falsehood" (1936/1946, 108). In contemporary jargon, he thought that they are not **truth-apt**, i.e. properly evaluated for truth or falsehood.

Fourth, he argued that ethical sentences do not express propositions but rather have only **emotive meaning** and not the sort of meaning normally associated with factual sentences.

As I'll explain below, these four striking tenets of Ayer's emotivism set the agenda for many of the criticisms of expressivism. In one way or another expressivists ever since Ayer have been trying to show how such conclusions are not forced by the basic expressivist idea—viz. that by focussing on what ethical statements do rather than on what they represent, we can come to appreciate that they express noncognitive attitudes towards the nonmoral world rather than describe moral properties of things, thereby avoiding ontological commitment to moral properties and facts.

Key Point: Emotivism is an early and extreme form of expressivism that conceives of ethical statements on the model of 'boo's and 'hooray's.

Disagreement and prescriptivism

Consider disagreement first. Whatever Ayer thought, intuitively speaking, it surely seems that there are genuine ethical disagreements. We are all familiar with disagreeing with someone about whether an action is wrong. Just think of the various opinions your friends and family may have about the ethics of Manning's leak of classified documents. And there seem to be deep-seated cultural differences that lead to some people favoring certain policies and laws that others oppose. To see this, just watch pundits discuss the pros and cons of a national referendum, or listen to a political debate in the lead up to a national election.

Of course, Ayer might put this appearance of moral disagreement down to disagreement about some underlying factual issue. For instance, consider two people who disagree about the value of same-sex marriage. Let's imagine the first person bases her objection to same-sex marriage on her belief that same-sex marriage decreases respect for marriage in society, leading to increased divorces and decreased well-being of children. The other person might think that if that were true, then same-sex marriage would indeed be bad. But let's imagine she doesn't agree about the factual issue about whether same-sex marriage leads to increased divorces and decreased well-being of children. If the apparently moral disagreement between these two people about same-sex marriage is based primarily on this difference of opinion about the consequences of same-sex marriage, then Ayer's description of the case looks pretty good. Their "moral disagreement" is really a disagreement about some underlying factual issue.

QU3: Which of the following are topics of factual disagreement that might underlie differences in moral opinion: (i) Whether certain symbols such as the Confederate flag or the swastika are disrespectful, (ii) Whether universal access to primary healthcare leads to lower overall healthcare costs than private insurance systems, (iii) Whether the death penalty is an effective deterrent of serious crime?

However, if you've ever been in an argument with someone who disagrees with you about same-sex marriage or any other heated ethical topic, you will surely have had the impression that even after clearing up all factual disagreements you'd still disagree with the other person. As we sometimes say: some people seem to have fundamentally different values from us. But, if that's right, then it seems that there can be genuinely *moral* disagreements, which is a serious problem for any view that implies—as Ayer's emotivism seems to—that there cannot be.

The main response to the objection about disagreement among expressivists is to back away from Ayer's suggestion that ethical statements express *mere* feelings—attitudes that two people may fail to share but which they cannot sensibly disagree

about—by arguing that there are other ways to disagree than two people having inconsistent beliefs. Stevenson already appealed to the idea of a “disagreement in attitudes.” For example, if you prefer we go out for dinner and I prefer we stay at home for dinner, then there’s a sense in which we disagree in preferences: we can’t both get what we prefer. (It’s important that these are preferences for what *we* do.) Similarly, imagine that your mother tells you, “Buy your miserable aunt a birthday present!” and your father says, “No, don’t do that! She’ll hate it no matter what.” Surely your mother and father have disagreed in what they told you to do. It’s impossible to follow both of their commands. Here we have a disagreement in prescription rather than a disagreement in belief.

Appreciating the relevance of prescriptions to explaining ethical thought and discourse, R. M. Hare (1952) defended **prescriptivism** as a theory of ethical language. The basic idea is that ethical statements are like commands in being prescriptive rather than descriptive, but they are unlike typical commands in embodying a “universalized prescription,” i.e. a sort of prescription that is supposed to go for everyone in similar circumstances.

With a much-discussed example, the **missionary and cannibals example**, Hare argued that expressivists not only have an explanation of genuine moral disagreement but this explanation is better than representationalist theories. In the example, he asks us to imagine that a missionary has landed on a distant island, equipped with a dictionary translating the local language into English. Observing the things they call “good” (in their language), the missionary notices two things. First, they seem to approve of the things they call “good,” striving for and promoting them. Second, the class of things they call “good” is radically different than the class of things the missionary would call “good”: the cannibals think people who are bold and collect a large number of scalps are good, whereas he thinks meek and gentle people are good.

Hare argues that if we thought that the word ‘good’ describes things as having a property, we’d have to conclude that the missionary’s word ‘good’ and the cannibal’s word for ‘good’ simply have different meanings. It’s like when British people use the word ‘football’ for a game played with a round ball and mostly the feet, and American people use this word to refer to a game played with an oblong ball mostly carried and thrown with the hands. If they don’t apply the word to the same things, that’s not because they disagree; rather it’s because they are using the words with different meanings.

But Hare thought it was pretty clear that the missionary and cannibals *do* disagree in their use of ‘good.’ Because of this, he thought that their disagreement couldn’t be about whether being bold and collecting lots of scalps has the property of being good but rather must be a disagreement in attitude and prescription. It’s not exactly clear what state of mind is expressed by a prescription, but it’s probably some desire-like attitude such as a preference. Hence, we might say that the cannibals prefer that people be bold and collect scalps, and using ‘good’ to express this preference they

make (on Hare's view) a universalized prescription to act boldly and collect lots of scalps; the missionary prefers something incompatible with this, and using the term 'good' to express this preference he makes (again, according to Hare) a universalized prescription to do something different.

If that seems like a plausible analysis of the example to you, then you might be inclined to agree with Hare that ethical language is prescriptive, and genuinely *moral* disagreement is a disagreement in attitude/prescription rather than in belief. However, others have worried that, although ethics and prescription seem intimately related, the prescriptive does not exhaust the ethical. For example, consider the idea that it'd be good if there were less people in the world (as that'd decrease the stress humans put on the natural environment). This seems like an ethical evaluation, but it is not obvious that it is making any particular prescription. (It surely doesn't prescribe individual people not to have children.)

Blackburn and Gibbard on taking ethics seriously

To see a way forward for expressivists, let's turn now to Ayer's idea that there's not much sensible work to be done in moral philosophy because moral language is really the expression of feelings rather than a way to report facts. When he made it, you can imagine that moral philosophers weren't too impressed with this suggestion! And there appears to be no reason the expressivist has to endorse it.

Ever since Hume, we've had the idea that some of our thought and discourse may initially seem as if it is representing reality even though it really involves some sort of projection onto reality of stuff that's actually in our minds. In Hume's colorful terms, we "gild and stain" reality with our sentiments.

To see what he has in mind, consider first a nonmoral example: thought and discourse about what's *delicious*. At first it might seem that a cone of Belgian chocolate ice-cream has the following two properties (among others): it is frozen and it is delicious. However, upon reflection we should start to see a difference here. The frozenness of the ice-cream is a product of the microphysical interactions of its atomic parts, something that would remain were no sentient beings to be around to consider it. By contrast, its deliciousness seems to be, at least in part, something we project onto the ice-cream. Indeed, it's not hard to imagine someone who finds Belgian chocolate ice-cream disgusting rather than delicious (perhaps they once ate enough of it to make themselves royally sick). Is it *really* delicious or disgusting? Probably neither, "really." This is different from the question of whether it is *really* frozen or liquid. This difference may hold, of course, without undermining our concern for figuring out which things are delicious and which are disgusting. After all, the fact that their deliciousness/disgustingness is in part "gilded and stained" by us doesn't change the fact that we find them delicious/disgusting, which involves taking pleasure/pain in consuming them.

QU4: Which of the following are plausible candidates for features we project onto reality as opposed to find there independently of our tendency to “gild and stain”: (i) The humor in a joke, (ii) The distance between New York and Sydney, (iii) The temperature of the earth’s atmosphere, (iv) The beauty in modern art?

There are two ways this idea has been used by expressivists to develop a version of the view that rejects Ayer’s contention that moral philosophy is mostly empty. First, Blackburn (1984) argued for a **projectivist** form of expressivism, according to which our ethical evaluations involve a kind of projection of our values onto the world around us. In his view, this actually aids the function of ethical thought and discourse. If, rather than seeing ethics as mere booing and hooraying, we see it instead as “gilding and staining” reality for the purpose of our own practical deliberation but also to influence others, we can start to understand how ethical discourse works to influence action and preferences within our moral community. Second, Gibbard (1990) argued for a version of expressivism he called **norm-expressivism**. One of the key ideas of this was that in making ethical statements we are expressing not a mere feeling but rather our commitment to systems of norms or general rules. His suggestion, then, was that these provide guidance on how to act when they are combined with specific beliefs about how reality is. (This is an answer to the question raised above about what specific desire-like state of mind is expressed by ethical statements.) Because of the social utility of being committed to similar rules for living together, Gibbard argued, there is good evolutionary reason to think that humans would have developed a way to voice their commitment to norms, as this provides a means for coordinating on basic rules for behavior. And this seems to be very important for finding a way to live together and improving it by continued dialogue and collective practical deliberation.

So, what we see in early work by both Blackburn and Gibbard is the possibility of an expressivist-friendly explanation of why moral philosophy is very important. This is not because it is aiming at a correct representation of some objective piece of reality. Ethics is different from physics. Rather, it’s because ethical discussion is a place where we express our values and commitments in order to coordinate and refine ways of living together. Moreover, as long as we think of these values and commitments as something meant for groups of people rather than individuals, Blackburn and Gibbard also develop resources to handle the challenge of disagreement discussed above. However, their idea about what ethical statements express leaves some important questions about ethical language unanswered. We turn to these next.

Truth-aptness and quasi-realism

While it may be a common thought that there is no objective truth in ethics only opinion, Ayer’s suggestion that ethical sentences are not truth-apt raises certain

problems in the philosophy of language having to do with the way we use the word ‘true’ and its relatives in ordinary discourse. Just consider the following mini-dialogs:

- (A) Lying is wrong.
- (B) That’s true, but this is a case of not telling the whole truth, not lying.
- (C) Torture is okay in extreme circumstances.
- (D) No, what you believe about torture is false. In fact, torture is always wrong.

These are ordinary things to say. But, if Ayer is right that ethical sentences aren’t even capable of being true or false, then how can it make sense for B to say “That’s *true*”? And if ethical statements express noncognitive attitudes rather than beliefs, then shouldn’t it seem nonsensical what D says about what C *believes* about torture? Moreover, if there are no ethical facts, how could torture in fact be always wrong?

Even if you’re inclined to say that ethical thought and discourse “gilds and stains” reality rather than represents it, you should worry about expressivism’s ability to make sense of these elements of ordinary ethical discourse. Using ‘true’ and ‘believes’ and ‘fact’ in conjunction with ethical statements doesn’t seem like nonsense. Moreover, as we’ll see in the following chapter, one of the main antirealist competitors to expressivism is the error theory, which says that our ordinary ethical discourse embodies an ontological error. So, unless expressivists can legitimate the use of these and related terms within ethics, it will seem that they avoid claiming that ordinary use of ethical terms is in error only at the cost of committing to the claim that ordinary use of terms like ‘true,’ ‘believes,’ and ‘fact’ is in error.

QU5: What does it mean to say that ethical sentences are truth-apt?

Blackburn (1984, 1993) pioneered the response to this worry that has become standard among expressivists. He argued that the difficulty is not with the expressivist’s view about ethical terms but rather with a common assumption philosophers make about ‘true,’ ‘believes,’ ‘fact,’ and related terms. The common assumption is that they apply only to areas of discourse and thought that are representational. This assumption stems from the **correspondence theory of truth**, which says that a proposition is true just in case it corresponds with reality. Importantly for Blackburn’s purposes, that is a highly controversial view about truth, and a competing view is the **minimalist theory of truth**, which says that ‘is true’ is a kind of linguistic or logical device which allows us to talk about (possible) statements or belief to which we assent. And since we can assent to ethical statements just as much as we can assent to descriptive statements, if one adopts minimalism about truth, there’s no reason the expressivist has to deny that ethical statements are truth-apt. Similarly, we might have thought that the word ‘believes’ refers to a mental state that represents reality as being some way. But if we instead think of it as referring to *whatever* mental state is involved

in assenting to a statement, then there will be no reason expressivists have to say that the mental states expressed by ethical statements aren't beliefs. Moreover, there seems to be little difference between believing or saying, e.g., that it's *true* that torture is wrong, and believing or saying it's a fact that torture is wrong. So minimalism about 'truth' and 'belief' seems to support minimalism about 'fact.'

Hence, by adopting minimalism about 'truth' and 'belief' and 'fact' (which admittedly are controversial views in their own right), it looks as if expressivists can avoid the worry generated by our mini-dialogs above. Blackburn referred to this line of thought as the beginning of **quasi-realism**, the program of winning back for the expressivist right to use terms, such as 'true,' 'believes,' and 'fact' that were previously thought to be the province of only realists. The resulting expressivist view is *quasi*-realist in freely saying that ethical statements are truth-apt and that they express ethical beliefs which may in fact be true, but then explaining why that doesn't commit the expressivist to realism.

Key Point: A minimalist theory of truth lets the *quasi-realist* expressivist say that ethical statements are truth-apt without undermining his commitment to the view that ethical statements do not represent reality.

Even if one goes this far with expressivists, one might worry that the view still entails that what's right and wrong depends on what we happen to think is right and wrong. As we'll see in Chapter 8, there are some metaethicists who accept some form of this mind-dependence claim about ethics. However, part of the continuation of the quasi-realist program for expressivists like Blackburn is to argue that once they have earned the right to use terms, such as 'true,' 'believes,' and 'fact' within ethical discourse, there's nothing to prevent them from denying claims of mind dependence. Consider, for example, the following mini-dialog:

- (A) Torturing people is wrong.
- (B) Sure, but it wouldn't be wrong if we all didn't disapprove of it.
- (A) No, that's not right; even if we all failed to disapprove of torture, it'd still be wrong.

You may agree with A or B, but the present point is that expressivism as a metaethical view is neutral on this issue, and most expressivists tend to agree with A. How could they do that if they think A's first statement is primarily an expression of an evaluative attitude? It's because they'll insist that A's second statement is *also* the expression of an evaluative attitude. It's something like the expression of a negative evaluation of the possible state of affairs where we all fail to disapprove of torture.

As I mentioned in Chapter 1, it is not entirely clear how we should characterize realism. There, I proposed to use, as a first approximation, a definition of it as the view that ethical facts obtain objectively. I mentioned (in note 1), however, that this definition

is imprecise because some philosophers might say, e.g., it's an objective fact that torture is wrong, but go on to deny that in saying this they are metaphysically characterizing what reality is like. The philosophers I had in mind are quasi-realist expressivists such as Blackburn and Gibbard. Their minimalist positions about 'truth,' 'belief,' and 'fact' put them in a position to say things about ethical facts without obviously committing to ethical facts as part of a full and final metaphysical characterization of what reality is like. If we want to continue with tradition and characterize their view as a form of ethical anti-realism, we'll need to define moral realism as the view that a full and complete metaphysical characterization of reality will include objectively obtaining ethical facts.

However, there are two difficulties about quasi-realism we might continue to worry about. First, if we "go minimalist" about 'truth,' 'belief,' and 'fact,' where do we stop? These concepts seem tied up with concepts like reference, property, representation, obtaining, objectivity, etc. And if we adopt minimalism about all of these terms, we might have no stable way to draw the distinction between expressivism and realism. Second, it's not even clear this "go minimalist" move works for all elements of ordinary ethical discourse which might have tempted us to think of it as representational. For instance, consider the term 'knows'; it's perfectly ordinary to say something like, "I know that I ought to conserve energy, but I always forget to turn off the lights." But how could one claim to have ethical *knowledge* if the expressivist is right that one's judgment that she ought to conserve energy isn't a representation of reality? Regardless of one's stance on minimalism, most views of propositional knowledge require that there be some fact that-*p* that is tracked by one said to know that-*p*. But the expressivist doesn't seem to be in a position to appeal to a tracking conception of ethical knowledge, since he's an antirealist about ethics.

Frege–Geach

Let's turn now to the most famous challenge to expressivism. It is widely assumed in the philosophy of language that the meaning of a whole sentence is a function of the meaning of its parts and the way they are put together. This **compositionality assumption** seems to offer the only plausible framework for explaining the amazingly quick speed with which humans learn their first language and how we understand indefinitely many novel sentences with ease. If this is right, then there must be some stable content that is contributed by any part of a sentence to the meaning of all of the sentences in which it figures.

For example, consider the sentences "Roses are red" and "Fire engines are red." Both contain the word 'red.' So, according to the compositionality assumption, there's likely to be something common to both occurrences—the meaning of 'red'—which this word contributes *systematically* to the meaning of the whole sentences in which it occurs. Likewise with the other words in these sentences (and any other sentence we can understand). If that's right, then we can think of language learners as cottoning

on to a finite number of word meanings and a finite number of principles for combining these meanings into meaningful whole sentences. And, with these two abilities, they should be able to understand an indefinite number of sentences.

Unfortunately, that much of modern semantics already stands in tension with early versions of expressivism, such as Ayer's emotivism. This was first noted by Searle (1962) and Geach (1965) and it rests on a point made by Frege; it has come to be called the **Frege-Geach problem**. The point made by Frege was that simple sentences could occur in more complex contexts, such as under negation ("It's not the case that..."), in a question ("Is it the case that...?"), under a propositional attitude report ("Sarah believes/doubts that..."), or in the antecedent or consequent of a conditional ("If..., then...").

To see how the objection goes, grant for the sake of argument that the following simple ethical sentences have no factual meaning and, as Ayer thought, only express negative moral sentiment:

"Tormenting cats is wrong."

"Getting little brother to torment cats is wrong."

QU6: How might we translate these two statements into the language of 'boo's and 'hooray's?

Then ask yourself what the emotivist would say about the meaning of more complex sentences which embed those sentences as parts. For example, consider the following conditional:

"If tormenting cats is wrong, then getting little brother to torment cats is wrong."

It doesn't seem that the emotive "meanings" of the simple sentences can be the content these sentences contribute to the more complex conditional sentence. For one needn't have any negative moral sentiment at all towards tormenting cats in order to endorse the conditional sentence. This is because one who asserts a conditional—if *p*, then *q*—can remain neutral about each of its components. For example, one could consistently say, "If tormenting cats is wrong, then getting little brother to torment cats is wrong," but "Tormenting cats is not wrong."

Geach pressed this point by arguing that expressivists make a **fallacy of equivocation** out of what appears to be a clearly valid inference. The clearly valid inference is this one:

- (1) Tormenting cats is wrong.
- (2) If tormenting cats is wrong, then getting little brother to torment cats is wrong.
- (3) *Hence*, getting little brother to torment cats is wrong.

His idea was that expressivists hold that the simple sentence in premise 1 has emotive meaning by virtue of expressing a negative moral sentiment. But, as we have just seen, they must hold that the antecedent of the conditional sentence in premise 2 has some other kind of meaning. This implies that one who infers from these sentences to the conclusion 3 has equivocated in the course of the argument. But that's clearly not what's going on in this argument.

Because of this way of pressing the point, many expressivist responses to the Frege–Geach problem have sought to vindicate the inference and ones similar to it. For example, some have followed Hare in the idea that the ethical sentences convey prescriptions of some sort and noted that there are independently motivated logical properties of prescriptions. Perhaps someone who asserts 1 and 2 is committed to the prescription contained in 3. (However, it's not entirely obvious what this prescription is.) Others such as early Blackburn have suggested that conditionalized ethical sentences are used to express a higher-order attitude towards the simultaneous possession of two lower-order attitudes. In this way we may see premise 2 as expressing disapproval of anyone who would have the negative sentiment expressed by 1 but not have the negative sentiment expressed by 3. (However, there seems to be a difference between disapproving of combinations of attitudes and thinking an argument is logically valid.) And still others such as Gibbard have developed more complex expressivist accounts on which the content of ethical sentences can be represented as combinations of factual commitments and planning commitments, where the logical constraints on planning help to explain the validity of the inference Geach considered. And more recently Blackburn has developed another view according to which the logical connectives help to keep track of the way various statements commit us to other statements, sometimes “tying us to a tree” of other commitments, whether they are evaluative or factual.

It would take a chapter of its own to discuss the ins and outs of these proposals. However, I want to close this section by suggesting that the validity of the inference is not the core problem. The core problem is how expressivists can develop an account of the content of sentences that treats an ethical sentence as having the same content whether it is asserted or embedded in a more complex linguistic construction. Just consider the way we could embed the simple ethical sentences above under the predicates, “It is doubtful that...” or “It might be the case that...”. Because we can also embed any other declarative sentence in these contexts, the compositionality assumption puts strong pressure on us to think that there is some uniform account of the meanings of the sentences that embed in these ways. The standard account in the philosophy of language has been that declarative sentences (both ethical and non-ethical) express *propositions*, and it is these propositions that remain the same across simple and embedded contexts. As a result, much recent work on expressivism has been about whether it can countenance such propositions for ethical sentences without giving up on the contrast it seeks to draw between ethical sentences and sentences that represent reality; and, if not, whether it can offer some general alternative that applies in both the ethical and non-ethical cases.

Motivational internalism

One last kind of objection I want to mention stems from the motivational internalist view that ethical thoughts can motivate action (though not infallibly) all on their own. As I mentioned above, this is a plank in one of the core arguments in favor of expressivism. Because of this, however, some have worried that challenges to motivational internalism are *prima facie* problems for the expressivist. For example, throughout history philosophers have puzzled over the phenomenon of **akrasia**, wondering how it is possible that one could think one ought to ϕ but when the time comes fail to ϕ . Indeed, even strong forms of this seem possible: where one thinks one ought to ϕ , forms a fully committed intention to ϕ , but when the time for ϕ -ing comes one's will turns weak; without changing one's mind about the ethical issue one nevertheless has no motivation to ϕ . There's naturally a lot of debate about whether this is really possible, but if it is, then it may seem to put pressure on the expressivist's view that these ethical thoughts are desire-like attitudes. For desire-like attitudes are simply the sorts of attitudes one needs in one's psychology to be motivated; so how could one simultaneously have them (in thinking that one ought to ϕ) but not have them (in suffering from *akrasia*)?

In response, expressivists might challenge the possibility of such an extreme akratic situation, arguing that even if one's will is weak, that's not because one has *absolutely* no motivation to perform the relevant action. It's rather because one's motivation (i.e. the attitude expressed by the ethical statement) is outweighed or undercut by countervailing motivations. Still some have worried that we can imagine someone who is seriously depressed, so much so that they have lost all evaluative affect. Nonetheless, remembering their previous ethical views, such a person might still make ethical statements. And the expressivist's view of these statements is that they function to express evaluative attitudes. Should we say that the seriously depressed cannot really make an ethical statement? It is not clear the expressivist *has* to say that, but this sort of case does pose a further challenge to expressivism. In evaluating how convincing you find it, you'll want to think about what it means when expressivists claim that ethical statements "express" evaluative attitudes.

• CONCLUSION

Expressivism in metaethics is a tradition that seeks to replace the representationalist and cognitivist assumption common to most other metaethical viewpoints with the idea at the interface of the philosophies of language and mind that ethical sentences express nonrepresentational states of mind. There's no agreement about what this state of mind is, but it's supposed to be the sort of state of mind that might explain the distinctive motivational or action-guiding character of ethical discourse. Some candidates are emotive reactions, preferences, desires, or plans. These show up in Ayer's emotivism, Hare's prescriptivism, Blackburn's projectivism and quasi-realism, Gibbard's norm-expressivism, and a number of more recent views developed within the tradition.

In any case, the view can be motivated in several different (and compatible ways) drawing on:

- Hume's Law and open-question arguments combined with naturalism;
- motivational internalism combined with the Humean theory of motivation;
- supervenience.

Yet expressivism faces a number of serious objections, including objections about:

- disagreement
- truth-aptness
- seriousness of ethics
- semantics
- mind-independence
- motivational internalism.

These objections have forced expressivists to refine their view in several significant ways. The jury is still out on whether the refinements meet the challenges in a way that maintains the core motivations for expressivism, especially on whether the compositionality assumption at the heart of the Frege–Geach problem can be respected by expressivists.

• CHAPTER SUMMARY

- Expressivists typically encourage us to begin our metaethical enquiry by asking not about (i) the nature of some kind of supposedly real thing: ethical value, but about (ii) the nature of one of our practices: making ethical evaluations. Then, their answer to question (ii) typically claims some interesting disanalogy between ethical evaluations and representations of reality.
- If you're attracted to the naturalistic worldview, expressivism can be motivated by the thought that ethics is different than description of the natural world (this is expressed in different ways in Hume's Law and the open-question argument).
- If you're attracted to motivational internalism about ethical judgments, then expressivism can be motivated by arguing that ethical judgments play a distinctive motivational role in a Humean psychology of motivation, suggesting that they must be "desire-like" rather than "belief-like."
- The apparent supervenience of the ethical on the natural provides one further way to argue for expressivism.
- Ayer's emotivism was an early predecessor to modern expressivism; he held that ethical statements have no factual meaning and only express positive or negative sentiments, which he took to mean that they are not truth-apt and that moral philosophy and ethical disagreement is mostly empty.
- Hare argued for a more nuanced prescriptivist view which had the resources to explain some moral disagreement as disagreement in prescriptions.
- Blackburn and Gibbard have argued for even more sophisticated forms of expressivism that make sense of the seriousness of ethics and the truth-aptness of sentences that express attitudes rather than representations.
- The most famous objection to expressivism is the Frege–Geach objection. Blackburn, Gibbard, and others have developed proposals for meeting it.

• STUDY QUESTIONS

- 1 If the open-question argument is an argument for nonnaturalism, then how could it also support expressivism?
- 2 Can you think of any counter-examples to motivational internalism about ethical judgments? How do these affect the case for expressivism?
- 3 Do the missionary and cannibals in Hare's example mean the same thing by their words 'good'?
- 4 What is "Frege's Point" in the Frege–Geach problem for expressivists? Why does it cause a problem for expressivists?
- 5 Do expressivists agree that ethical issues are important issues and that there can be true ethical judgments?
- 6 Explain supervenience.
- 7 What's the difference between expressivism and the subjectivist view that ethical statements are about what their author likes and dislikes?

• FURTHER RESOURCES

- Bar-On, Dorit, and Sias, James. 2013. "Varieties of Expressivism," *Philosophy Compass*, 8 (8): 699–713. [Covers more forms of expressivism, also as applied in domains besides ethics.]
- Chrisman, Matthew. 2011. "Ethical Expressivism," in *The Continuum Companion to Ethics*. Continuum. [Covers more forms of expressivism.]
- Glanzberg, Michael. 2014. "Truth," *The Stanford Encyclopedia of Philosophy* (Fall 2014 Edition), Edward N. Zalta (ed.), <http://plato.stanford.edu/archives/fall2014/entries/truth/>. [Comprehensive discussion of many different theories of truth, including correspondence theory and minimalism.]
- Schroeder, Mark. 2010. *Noncognitivism in Ethics*. Routledge. [Provides a book-length introduction to noncognitivist views which include contemporary versions of expressivism.]

• ANSWERS TO QUESTIONS OF UNDERSTANDING

- QU1: In Chapter 1 we appealed to the idea of "different directions of fit" with the world to distinguish descriptive attitudes from directive attitudes. The idea was that beliefs try to fit the way the world is (in order to be true) whereas intentions, desires, plans, etc. try to make the world (come to) fit them (in order to execute or satisfy them). Maybe there are some nonrepresentational attitudes that are not directive, but the main nonrepresentational attitudes we'll be concerned with here are directive rather than descriptive.
- QU2: Describing yourself as feeling some way (e.g., as sad) expresses a belief about the way the world is, it's just that this particular way the world is involves

you and your feelings. This description could also be given third-personally by someone else observing you. By contrast, expressing your feeling (e.g., your sadness) conveys it to your audience. This is something you could do with a facial gesture, a cry, or an exclamation. It is somewhat controversial what to say about explicit avows such as “I’m feeling so sad!”—do they describe oneself as feeling sad or do they express the sadness, or perhaps both?

- QU3: Issues (ii) and (iii) are plausibly factual issues, whereas issue (i) involves a moral issue about what constitutes “disrespect.”
- QU4: Features (i) and (iv) are the most plausible candidates for features which we project onto reality rather than find there, although note that some philosophers would argue that even (ii) and (iii) are in some sense mind-dependent.
- QU5: It means that they are normally either true or false. Notice how the sentence “It’s sunny today” is something that we can normally evaluate as being either true or false, but “Take some sunscreen!” is not something we can really claim to be true or false. (In case you’re wondering, the “normally” is a hedge to allow for cases of vagueness and presupposition failure where some philosophers want to argue that the relevant statement is neither true nor false, but for different reasons and in a different way from the emotivist’s claim that ethical statements are neither true nor false.)
- QU6: We’d say something like “Boo for tormenting cats!” and “Boo for getting little brother to torment cats!” Notice how the if–then sentence which follows next in the text is not so easily translated into the language of ‘boo’s and ‘hooray’s.

• WORKS CITED

- Ayer, A. J. 1936/1946. *Language, Truth and Logic*. 2nd edn. London: V. Gollancz Ltd.
- Blackburn, Simon. 1971. “Moral Realism.” In *Morality and Moral Reasoning*, edited by J. Casey. London: Methuen.
- . 1984. *Spreading the Word: Groundings in the Philosophy of Language*. New York: Oxford University Press.
- . 1993. *Essays in Quasi-Realism*. New York: Oxford University Press.
- Geach, P. T. 1965. “Assertion.” *The Philosophical Review* 74 (4): 449–65.
- Gibbard, Allan. 1990. *Wise Choices, Apt Feelings: A Theory of Normative Judgment*. Cambridge, MA: Harvard University Press.
- Hare, R. M. 1952. *The Language of Morals*. Oxford: Oxford University Press.
- Searle, John R. 1962. “Meaning and Speech Acts.” *The Philosophical Review* 71 (4): 423–32.
- Stevenson, Charles L. 1937. “Emotive Meaning of Moral Terms.” *Mind* 46: 14–31.