

5

• naturalism

All of the views we have discussed so far¹ accept the following claim:

(C) If reality includes ethical facts, then they are not purely natural facts.

Nonnaturalists move from left to right: accepting that reality includes ethical facts, they try to figure out how to characterize them as nonnatural. By contrast, expressivists, error-theorists, and fictionalists move from right to left: being skeptical of facts that cannot be characterized in purely natural terms, they try to make sense of the idea that a full and complete metaphysical characterization of reality might in the end not mention ethical facts. But what if (C) is false? In Chapter 2, we canvassed arguments for thinking that ethical facts, were they to really obtain, would have to be nonnatural; and you'll recall that these arguments aren't airtight. So, perhaps we should consider the possibility that reality includes ethical facts *and they are natural*.

Ethical naturalism is the view that reality includes ethical facts and they are reducible to or otherwise can be made to fit within the facts we regard as natural. Different versions of ethical naturalism are distinguished by different strategies for explaining how ethical facts might in this sense be natural. Many philosophers think it would be nice if one of these strategies work, since we could then embrace a cognitivist conception of ethical thought in our philosophy of mind that makes it easy to understand why we commonly speak of ethical beliefs and ethical knowledge. We could also embrace a representationalist conception of ethical language, which coheres well with the way ethical sentences seem to work semantically more-or-less like ordinary descriptive sentences. Moreover, we wouldn't have to secure these views by recognizing nonnatural facts in our metaphysics. Indeed, ethical naturalism holds out the prospect of conceiving of ethical knowledge as just one more species of our knowledge of the natural world.

The key for any form of ethical naturalism is to develop a plausible explanation of why ethical facts are a species of *natural* facts. Unfortunately, as we discovered in Chapter 2, there is no uncontroversial characterization of what it takes for something to be *natural*. Nevertheless, nonnaturalists argued that ethical facts don't fit with other recognizably natural facts, such as physical, biological, psychological, etc.

facts in our conception of reality. So the main challenge for ethical naturalists is to convince us that this claim is false. In this chapter, we'll consider four strategies for doing so. The first challenges our previous assumptions about what counts as *natural*. The second challenges our previous assumptions about what it takes for there to be facts obtaining *objectively*. And the third and fourth argue for a new way of *reducing* ethical facts to facts discoverable by empirical science such as physics and chemistry.

Key Point: Metaethical naturalists are realists who think ethical facts are not “autonomous” from other natural facts but can be reduced to or otherwise made to fit with the sorts of facts we regard as making up the natural world.

• NEO-ARISTOTELIAN NATURALISM

Although we haven't been able to define what it is for a fact to count as natural, we have been working with the assumption that these are the sorts of facts discoverable by scientific methods as exemplified most clearly in modern physics and chemistry. However, there is a tradition stemming from Aristotle that places considerable weight on the idea that there is a further sort of natural fact, one which often requires significant metaphysical reflection to discover. This further sort is facts about things' *natures*. For example, we speak about “human nature.” We also speak about “what's natural for x ,” where x picks out some species or natural grouping of beings; e.g. “Rough skin is natural for the common toad.”²

The important idea in the present context is that these *natures* so-conceived are supposed to be *natural*, which means that facts about them are natural facts. If that's right and it's possible to derive significant ethical conclusions from facts about human nature, then ethical facts will begin to look like a species of natural fact. They won't be “natural” in the sense of discoverable by recognizably natural sciences such as physics and chemistry. But they'll still count as “natural” under a somewhat broader conception of the natural.

So, let's accept for the moment that facts about things' natures are natural facts. How do we get from there to locating ethical facts within the natural world? There are two related ideas that help to make this connection.

The first idea comes from Geach (1956), who famously distinguished between **attributive adjectives** and **predicative adjectives**. The distinction is best understood through the test he proposed. Geach asked us to consider phrases of the form “is an A B,” where A is an adjective such as ‘red’ or ‘big’ and B is a noun like ‘book’ or ‘flea.’ When an instance of this formula can be split into “is A and is a B” without changing meaning, Geach refers to the use of A as *predicative*. The idea is that, in these cases, saying that something is an A B amounts to predicating A of it (and predicating B of it). By contrast, when you cannot split up the predicate in this way, Geach refers

to the use of A as *attributive*. The idea is that, in these cases, saying that something is an A B amounts to qualifying one's use of B by A. In his example, "is a red book" predicates being red (and being a book) and so is predicative. However, "is a big flea" does not predicate being big (full stop) but only being big compared to other fleas.

QU1: Which of the following are attributive uses of an adjective: (i) a flat desk, (ii) a harmful comment, (iii) a laughing child, (iv) a costly meal?

Now, to see how this is relevant to metaethics, consider our previous question about the meaning of 'good.' Following Moore, we asked "What does 'good' mean?" and considered answers such as that it refers to a simple and irreducible nonnatural property. That means that, in seeking to answer that question, we were tacitly assuming that 'good' is a predicative adjective. We were, that is, asking what is meant when we say that something is *good*. But Geach wondered: what if it doesn't really make sense to ask if something is good (full stop); what if 'good' is an attributive adjective rather than a predicative adjective? More precisely, Geach argues that a lot of metaethics done in the wake of Moore made precisely this mistake. His idea was that a lot of the debate about what it is for something to be good (full stop) is misguided much like debate about what it is for something to be big (full stop) would be. We can of course sensibly ask what it is for something to be big or good *for an x* (e.g., flea, city, negative number, galaxy, philosophy department...). But Geach thought it was nonsense to ask what it is for something to be (simply) good.

His positive suggestion was that metaethicists should focus instead on what it is for someone to be a good *person*, and this will have something to do with what makes someone a good instance of the kind *human being*. But how do we go about identifying the attributes that make something a good instance of its kind, e.g., a good human being? Here is the second important idea: we can identify these by appealing to functions or (what is very similar) characteristic ways of living.

To see what this means, consider first the kinds of things we can easily define by appeal to their function (what are sometimes called **functional kinds**). For example, a knife is (more or less, by definition) something that's for cutting. Famously, Aristotle argued that, for things with a function, we can identify properties that make them a good instance of their kind by thinking about what would help them to perform their function well. Of course, not just anything that can be used to cut is a knife (random shards of glass can also cut); but with a knife, we can say that having the function of cutting is part of what it is to be a knife—part of its "nature." So a good knife will be one that cuts well. (Contrast: a shard of glass may cut well, but that doesn't seem to make it a good shard of glass.)

QU2: Which of the following are functional kinds: (i) car, (ii) moon, (iii) lamp, (iv) human being?

Given this, we can then ask which property of a knife would help it to cut well? (Ignoring the complication that different knives are for cutting different kinds of things in different circumstances), in general, it's something like *being sharp*. Hence, we might say that sharpness is a property that makes a knife a *good* knife. And voilà, from considerations about something's nature, we have reached a conclusion about what makes it a *good* instance of its kind.

QU3: What are the features that make the following functional kinds good instances of their kinds: (i) shoes, (ii) power drill, (iii) textbook?

But there's a large gap between the nature of things like knives and *human* nature. Knives are pretty clearly for cutting, but what are humans for? It is not obvious how to answer this question. To see how we might close that gap, it's helpful to consider the natures of plants and animals first. It's controversial whether living things such as plants and animals have functions (things they are for), but we might view them as having something similar: **characteristic ways of living**, which can be more and less fully realized much like (though in a more complicated way than) knives can be more and less good for cutting.

It's controversial whether it makes complete sense that something's characteristic way of living generates standards for what counts as it living well. For it can be much harder to determine something's characteristic way of living. Aristotle's idea seems to have been inspired by the division of labor outlined in Plato's description of the ideal city, so another way to think about "characteristic ways of living" might be in terms of one's *job* or *role* in a group.

In any case, if we accept that living things have characteristic ways of living, then, following Aristotle, we seem to be in a position to identify the properties that make plants and animals good instances of their kind. These will be the properties that help them to realize the characteristic way of living for their kind of being. For example, oak trees characteristically grow to 65–130 feet. So the properties that make a particular plant a good oak tree are (among other things) the ones that help it to grow to 65–130 feet, e.g., a strong root system, leaves that are resistant to parasites, etc. (Note these properties are not identified as the properties that make a particular oak tree *good for* any particular human purpose, such as having a shaded yard, but rather as the properties that make it *good at* realizing the oak tree's characteristic way of living, or its "role" in the forest.)

Similarly for animals: the angelfish normally hides alongside a vertical rock face in order to avoid predators and ambush prey. So the properties that make a particular angelfish a good angelfish are (among other things) the ones that help it to hide alongside a vertical rock face, e.g., being thin and striped in colors like plants that grow up rocks. (Again, note these properties don't have anything specifically to do with human purposes for an angelfish, such as having beautiful fauna in one's aquarium; what matters are the properties that aid in achieving the organism's characteristic way of living.)

So we might say that the properties of being a good oak tree or being a good angelfish are constituted by a complex set of properties such as having a strong root system, having leaves that are resistant to parasites (in the first case) and being thin and striped in colors like plants that grow up rocks (in the second case). These, note, needn't be *statistically* common features in order to count as elements of the "characteristic" way of living. In some species, it may be only a select few who grow old enough to live in the characteristic way. Still, as long as we can reasonably identify a characteristic way of living for some living being, there will be properties that help it to realize this; and these will count as what makes something a good instance of the species.

The important point about these examples for metaethical naturalism is that it's plausible to think that these properties are *natural* properties; however, they might be thought to imply normative facts: how oak trees and angelfish *should* be. Going down this route promises a new way to show such 'should's to be about natural facts.

QU4: How would a Neo-Aristotelian derive facts about how dairy cows should be?

Surely, however, the should-facts about oak trees and angelfish are not *ethical* facts. **Neo-Aristotelians** such as Hursthouse (1999) and Foot (1999) have argued that, whether or not we're inclined to count facts about how plants and animals should be as ethical facts, we can extend the basic framework to humans and then use reflection on *human* nature to determine how humans should be. This extension will give us purchase on a kind of *normative* facts, which are natural. Then from this we might derive facts about what actions are right and wrong. And these will begin to look much more like ethical facts.

Initially, though, that seems wrong. Humans are, of course, animals of a particular species, and we can grant this means that there are ways humans should be, so to speak, *biologically*. For example, humans should *biologically* have fully opposable thumbs; having these is a feature that helps us to realize the characteristically human way of living. Maybe this even implies that there is a right way for humans to do things like hold a hammer. But these facts, which we can grant to be natural because biological, are hardly *ethical* facts. So, if all we get are biological should-facts about humans, the Neo-Aristotelians haven't yet shown ethical facts to be natural facts (or how to derive ethical conclusions from facts about our natures).

Neo-Aristotelians are aware of this challenge, and they develop different ways of addressing it, deriving from considerations such as the positions of humans in complex social networks, our ability to engage in reflective thought, and the capacity we have to feel particular sorts of emotions. The details are somewhat complicated, and it is beyond the scope of this book to consider which of these routes is the most promising way to develop a Neo-Aristotelian view.

However, I do want to mention one important idea developed by Hursthouse and consider some of the challenges it faces, in order to give you a sense of how debate

about a particular metaethical position might go. Hursthouse's idea is that *rationality* is an especially important feature of human nature. It gives us the ability to act from reasons rather than impulse. It also puts us in a position to reflectively endorse or disavow other features of our natures, in effect letting us choose the type of person we want to be. Accordingly, Hursthouse argues that some of the ways humans should be are special. She privileges those ways humans should be that derive from our capacity to be rational. While still belonging to the class of properties that help us to realize ways of living characteristic of human beings, these properties are not, in her view, merely biologically good-making features of a human being. They are *ethically* good-making features of a human being: what we might reasonably call **ethical virtues**. From the ethical virtues, Hursthouse proposes to derive facts about what actions are ethically right and wrong in various circumstances. To make a long story short, she does this by appeal to what the virtuous person would characteristically do in those circumstances.

Some will worry that the very specialness of rationality undermines any claim Hursthouse might have had to be an ethical *naturalist*. After all, we might grant that the ability to act from reasons distinguishes humans from other animals, but so do many other abilities and tendencies. From a biological, ethological, or even sociological point of view it is not clear why we should single out rationality as special in the characteristically human way of living. Indeed, some may even think that it's characteristically human (at least for our particular epoch) to be able to do something like embrace the absurdity and radical freedom that comes with a wholesale rejection of rationality. So, we might wonder: what makes our rationality so special?

There's one answer to this that Hursthouse cannot give without undermining her aspirations to defend a form of naturalism. She cannot say that ethics is *autonomous* in the way it is seated in our rationality and not in the other aspects of our natures. The problem with this answer is that it, in effect, assumes that one feature of our natures (and the should-facts that derive from it) is special, i.e. "sui generis" with respect to the other features that are recognizably natural. Hence, we'd be left without an answer to the question all naturalists have to answer: how is it that ethical facts reduce to or otherwise fit with recognizably natural facts? Hence, the nonnaturalist would see this manoeuvre as a concession that ethics is sui generis. I should note that Hursthouse does respond to this objection by, in effect, challenging the conception of something's being natural implicit in the objection. She thinks it is too narrow. This is why I earlier characterized the position as one that attempts to undermine the previous conception of what it takes for a fact to be natural.

Another worry we might have, however, is whether Neo-Aristotelians end up with a plausible account of what the ethical facts are. After all, there are many characteristically human traits, even some that might depend crucially on our rationality, that can seem to encourage morally questionable behavior. For instance, waging tactical wars, rearing animals for food, lying about sexual infidelity, favoring others simply because of their religion, and finding and exhausting fossil fuels are all parts of what

an ethologist might describe as characteristically human ways of life. But surely we can wonder whether those are ethically good ways for us to be. The problem here is one of methodology: if we're meant to give an account of ethical facts that makes them fit within the natural world, but the account we end up giving depicts ethical facts as radically different than we pretheoretically thought, we might begin to wonder whether the account is really an account of what we wanted to account for.

In response, a Neo-Aristotelian might try to focus on characteristically human ways of living that we pre-theoretically accept as ethically good, seeking to work towards reflective equilibrium between our pre-theoretic intuitions and a systematic ethical theory. This methodology is characteristic of much ethical theorizing. When it comes to identifying ethically good ways of living that are characteristically human, we face a significant risk that there will be so much variability across different cultures and subcultures that nothing will emerge (as Hursthouse claims rationality does) that can be used to define the specifically ethical virtues.

Key Point: Neo-Aristotelian naturalists seek to derive ethical conclusions from premises about human nature. If these premises can be reasonably treated as natural facts, the ethical facts that follow from them will have some claim to fit among other natural facts.

• RELATIVISM AS A FORM OF NATURALISM

Let's turn now to a different kind of metaethical naturalism. A common idea from anthropology is that right/wrong are *relative to* different "moralities."³ In one sense, this may be a purely anthropological claim: different groups of people adhere to different norms, and the norms they adhere to determine what kinds of things they think are morally right and wrong. (We already saw echoes of this idea in Hare's missionary and cannibals example from Chapter 3.) Such **anthropological relativism**, however, doesn't yet amount to a view about the *nature* of ethical facts. To get the view called **metaethical relativism**, we need to think further that there's no such thing as an action being right/wrong full stop but only relative to this or that morality. Compare: there is no such thing as being on the left/right full stop but only relative to some perspective or point from which one is looking at something. Sometimes the basic idea is put by saying that morality isn't *objective*. Although suggestive, it's worth heeding two warnings about this way of framing metaethical relativism.

First, sometimes when we say that something isn't objective, we mean that it's subjective. And there is an extreme form of relativism known as **subjectivism**, which says that there's no such thing as an action being right/wrong full stop but only relative to this or that person's own personal values. However, subjectivism is not the only kind of metaethical relativism. The kind inspired by anthropological relativism is one

which relativizes ethical facts to different *groups* of people. So we shouldn't think that metaethical relativists have to be subjectivists, though some are.

Second, the idea that morality isn't objective but rather relative problematizes the definition of ethical *realism* I gave in Chapter 2. There I defined realism in terms of a commitment to ethical facts obtaining objectively. But this raises the question: should we count the relativist as a realist or antirealist? Often metaethical relativists, and especially subjectivists, have been classified as antirealists, but depending on how we understand the term "objectively" there's a way to classify at least some of them as realists. In one sense, if some action is wrong only in the sense that it is forbidden by some person's or group of people's values, that means its ethical status is not *universal*. But the fact that some person's or group of people's values forbid the action would still seem to be a fact that is fully *out there* and not depending metaphysically on the "projection" or "construction" of human thought. To be sure, it may be a fact that is at least partly psychological in nature, but facts about a person's or group of people's psychologies might be reasonably thought to be perfectly natural facts.⁴

So, when some metaethicists deny that ethics is objective, it is not always clear what they mean (or what their critics mean when criticizing them for making such claims); and this interacts in a confusing way with the definition of realism. For the relativist's position to be a form of *naturalist* realism it must rest on the suggestion that ethical facts are "out there" but not *universal* in character. If we accept this, another kind of fact begins to look like a plausible candidate for being both ethical and natural: facts about the values or norms of different people or groups of people. Maybe, to pick up on the non-subjectivist forms, the facts about what is ethically right/wrong and good/bad are relative to different cultures very much like the facts about what is legal/illegal are relative to different jurisdictions or facts about what is polite/impolite are relative to different social groups. Is it wrong for couples to live together before they are married; is military valor good; should the dead be buried or burned? Maybe there are facts answering each question, but they depend tacitly on which moral community we are talking about.

There are many ways to develop this idea, and we won't be able to consider all of them here. But one characteristic relativist move is to argue that statements about what's morally right or wrong cannot be evaluated as true/false *simpliciter* very much like statements about what is large or small cannot be evaluated as true/false *simpliciter*. In both cases, relativists think we must fix some further **implicit parameter** in the statement before we can evaluate it as true or false. In the case of statements about size, pretty much everyone is a relativist in the sense that they agree that we need to determine a comparison class before we can evaluate the statement: When you say "Fluffy is small," do you mean small for a dog or small for a guinea pig? Before we know what you're comparing Fluffy to, we can't even begin to figure out whether what you say is true. Ethical relativists often argue something similar is true of statements about which actions are ethically right/wrong. The idea is that these depend crucially on the value of some further implicit parameter.

QU5: What are the implicit parameters in the following sentences: (i) The table is flat, (ii) It's a long way to Houston, (iii) Sarah is very old?

In this vein, Harman (1975) argues that statements about what's ethically right/wrong make sense only relative to what he calls "agreements." By this he means explicit or implicit intentions among a group of people to adhere to some schedule, plan, or set of principles, under the understanding that others also intend to adhere to it as well, under a similar understanding. To see what he has in mind, imagine for instance two different communities of people: Community A "agrees" to bury their dead relatives; community B "agrees" to cremate their dead relatives. Now consider this statement: "It's ethically wrong not to bury a relative who dies." Is this statement true or false? Harman's answer is that there's no *objective* (in the sense of "universal") fact of the matter; rather it should be considered true relative to community A's agreement and false relative to community B's agreement.

In order to be more precise, we might say that this kind of ethical relativist is an antirealist about *universal* ethical facts, but they are realists about *relativized* ethical facts. Of course, everyone should agree that there are different things demanded by different cultural norms, but the relativist thinks these non-universalized norms are what properly ethical facts are about.

Explaining how the moralities of different cultures come about is a fascinating and difficult topic for ethologists, anthropologists, psychologists, sociologists, and historians. But if we assume that what they are explaining is as much of a feature of the natural world as the other sorts of culturally variable facts they attempt to explain (e.g., facts about what's legal/illegal or polite/impolite), then there will be a good case for thinking of ethical facts (understood as the relativist does) as a species of natural facts. So, in this way, ethical relativism can be seen to be a form of naturalism.

Before moving on to a different kind of metaethical naturalism I want to consider some objections and replies so you can see how debate over metaethical relativism might go. We might worry, first, that there's a crucial difference between one's view about what's polite/impolite and one's view about what's morally right/wrong. (The case of what's legal/illegal is complicated by the fact that some philosophers think legal duties are a species of moral duties, so it's controversial whether legal opinions are more like moral opinions or more like opinions about etiquette in the following regard.) The difference is that one's moral opinions seem to reveal at least partially how one is disposed to behave under various circumstances. If you tell me that giving to charity is morally required for people as well off as us, then it's reasonable for me to expect you to be at least somewhat motivated to give to charity. If it looks like you are completely cold to the prospect, I'll begin to wonder whether you were lying to me or didn't understand what you were saying. The same does not seem to be true for etiquette. If you tell me that it's impolite to eat with one's fingers, I could easily make sense of your total lack of motivation to follow this rule. You don't need to

have lied or misunderstood what you were saying; maybe you just think the rules of etiquette in our culture (or at least this rule) are stupid. So, even if ethical statements are properly construed as representations of perfectly natural facts (roughly speaking, facts about what's right/wrong relative to different norms), they cannot be *mere* representations of these facts insofar as their connection to their author's motivations is markedly different from the connection born by statements about other culturally relative facts, such as etiquette.

This objection assumes motivational internalism, the doctrine we first encountered in Chapter 1, which says that an ethical judgment can motivate action without the help of any further desires, inclinations, or emotions. There is no temptation to think judgments of etiquette are like this, but there is some temptation to think ethical judgments are like this.

One response to this objection developed by Dreier (1990) is to argue that ethical statements are implicitly **indexical** in a characteristic first-personal way. Linguistically, indexicals are words such as 'here,' 'now,' and 'I,' which are used as indices referring in a systematic way to different things in different people's mouths. For example, if I say "I am happy," the word 'I' refers to me, whereas if you uttered the same sentence, the word 'I' in your mouth would refer to you. Generally speaking, 'I' is used to refer to the speaker him/herself. Some statements, it seems, may be *implicitly* indexical. For example, if I say "Tacos are yummy," the word 'yummy' means something like *generally tasty to me*, where this 'me' refers to *me*, the speaker of that sentence. In your mouth 'yummy' would mean something like *generally tasty for you*, where this 'you' refers to *you* the speaker of that sentence. I can of course make statements about what is yummy-for-you, but these won't connect to my culinary inclinations in the same way as my statements about what is yummy.

QU6: What are the implicit indexicals in each of the following sentences:
(i) She's in front, (ii) Yesterday it was sunny?

Defending a subjectivist view, Dreier argues that statements about what is good are like this. One can make statements about what is good-relative-to-Western-values or good-relative-to-capitalist-values without displaying any motivation to action. But when one makes a genuine *value* judgment about what is (simply) good, this should be understood as implicitly relative to one's own values. If part of what makes a statement about what's good a genuine value judgment is that it is implicitly indexed to the speaker's own values, then it would be unsurprising that we are generally though defeasibly motivated to act in accordance with our ethical statements.

A related objection to relativism is that ethical facts (if there are such things) seem capable of generating (justifying) reasons to act independent of our desires and concerns, whereas facts about etiquette don't seem to be like this. If it's morally wrong to steal money from the tip jar, that's a reason for you to refrain from doing this—no

matter whether you don't care about what's morally wrong or anything related to it. By contrast, the intuition is that the fact that it's impolite in some cultures to talk on the phone while using the toilet can be a reason to refrain from doing this, but only if you care about being polite (or at least something related to it). On this score, the relativist's ethical facts appear not to be reason-generating in the right sort of way. The fact that it's wrong relative to some "agreement" not to bury one's dead relatives doesn't give me a reason to do anything, unless of course I care about that agreement.

Key Point: Relativism is not always counted as a form of naturalist realism, but if relativists view ethical facts as non-objective because relative to something like different (people's or cultures') moralities, those facts can still be objective in the sense of "out there" for discovery by empirical science. In this case, it could be an attractive route for "locating" ethical facts among the natural facts.

Two further objections to relativism are worth mentioning. First, the standard objection to relativism is that it cannot make proper sense of moral disagreement. To be sure, Harman's relativism conceives of ethical facts as relativized to "agreements" we might have with the people with whom we live together in community. But what about the case where people from radically different cultures are discussing moral matters? Here, again, Hare's missionary and cannibal thought experiment from Chapter 3 is relevant. Part of the power of that example is that it really does seem like the missionary and the cannibals disagree about what is right/wrong. As far as relativism is concerned, however, unless they somehow share a morality, they are talking past one another. This suggests the second objection: relativism is based on the suggestion that there are different moral communities of people living by different moral rules and we're wrong to look for any universal rules. But what constitutes a moral community? The problem isn't that we don't have implicit agreements with many people about how to live together but rather that we have lots of these cross-cutting any plausible way of drawing the line between those who and those who do not share a moral community.

• A POSTERIORI NATURALISM

We turn now to some more complicated approaches to locating ethical facts among the natural facts. Although complicated, these turn on views about language and mind that have been worked out mostly independently of metaethics. So they might seem to be neutral tools that we can deploy from these other areas of philosophy to try to make progress in metaethics.

One prominent strategy for locating ethical facts among the natural facts depends on some ideas that gained currency in philosophy of language and philosophy of science

in the 1980s. You'll recall from our discussion of the open-question argument that Moore thought we can say things of the form "I know x is N , but is x good?" (for many natural properties N). From this he concluded that the property of goodness is not plausibly thought to be the same as any of these natural properties (e.g., the property of causing pleasure, maximizing happiness, etc.). In response, however, many naturalists have pointed out that most of the properties we think of as natural also wouldn't pass this test. The standard example is the property of being water. It is now quite widely known that water is H_2O , but before eighteenth-century chemists discovered this, someone could sensibly say "I know that x is H_2O , but is it water?" That is, they could ask this question without revealing confusion about the meaning of the relevant words or a less-than-adequate grasp of the relevant concepts. Indeed, someone could still sensibly ask a question like this today as long as they were simply ignorant of the chemical composition of water. This is because the chemical composition of water is something we discover **a posteriori**, i.e. via empirical investigation of the world rather than via **a priori** reflection on our concepts and the meanings of our words.

This means that the property F and the property G might be identical, even though our concept for F -ness is not identical to our concept for G -ness. Being water is identical to being H_2O , but one doesn't display conceptual or semantic confusion by wondering whether something that is H_2O is water; that just displays ignorance about a chemical fact. Similarly, some ethical naturalists have argued that ethical properties might be natural properties—just not in a way that would close Moore's open questions. Instead these naturalists propose to seek to fit ethical properties and facts within the natural world via a posteriori investigation of the world rather than via a priori analysis of our ethical concepts.

QU7: Why is the fact that water is H_2O something that cannot be discovered a priori?

It's one thing to point out this lacuna in the open-question argument, it's another thing to propose a plausible form of a posteriori naturalism, one we might imagine confirming scientifically like chemists confirmed the molecular theory of water. It's not like we can collect a bunch of goodness in a test tube and analyze its molecular structure. Accordingly, most a posteriori naturalists concede that ethical properties are going to be somewhat complex and more difficult to analyze than the property of being water. However, one cause for optimism stems from the **causal theory of reference**. The basic idea of this theory is that some words refer to individuals or properties not because of anything we might find via analysis contained in the concepts these words express but rather because our use of these words is causally tied to those individuals or properties.

Kripke (1980) argued that this is true of *names*. He used an example of the name 'Gödel,' with which we might associate various ideas such as *being the mathematician*

who proved the incompleteness theorem. Now imagine, however, that it turns out that we are mistaken, and someone else named ‘Schmidt’ was actually the person who proved the incompleteness theorem and Gödel stole his proof and became famous because of it. Kripke suggests (quite plausibly) that if this were the case we wouldn’t say that the name ‘Gödel’ actually refers to Schmidt in virtue of his being the one who really proved the incompleteness theorem. Rather we’d say that we were mistaken about Gödel’s mathematical accomplishments.

QU8: In Kripke’s thought experiment, who does the name ‘Gödel’ refer to?

From this, Kripke concludes that the referent of a name is determined not by what ideas we associate with the name but rather by whatever individual happens to be causally linked to our use of the name. The important point for our purposes is that figuring out which particular person the name ‘Gödel’ refers to requires a posteriori investigation of the causal chain between our use of this name and the person originally baptized ‘Gödel’—not a priori analysis of the ideas we associate with this name.

The reference of names is a controversial issue in the philosophy of language, but in a similar fashion, many philosophers of language following Putnam (1975a) have thought that natural kind terms, such as ‘water’ get their referents not from the ideas (e.g., wetness, clearness) we happen to associate with them but rather from causal chains between their regular use and the instantiation of some property (e.g., the connection between ‘water’ and instantiations of H_2O). The basic idea is that we need a posteriori investigation of the world rather than a priori analysis of concepts in order to determine what natural kind terms such as ‘water’ refer to.

If this is right, and we assume that ‘good’ is also a natural kind term, then figuring out the nature of goodness (just like figuring out the nature of water) would require a posteriori investigation of the causal chain between our use of this term and the property to which it refers. If this is right, then we shouldn’t expect a priori conceptual analysis to be probative with respect to the nature of goodness any more than we expect it to be probative with respect to the nature of water or with respect to the prover of the incompleteness theorem.

That’s a significant step towards a new a posteriori form of metaethical naturalism. However, even if we accept the causal theory of reference and apply it to ‘good,’ that still doesn’t tell us what goodness is, such that this property can be seen to reduce to or otherwise fit with natural properties. In response, a posteriori naturalists tend to say that goodness is probably going to be quite complicated, but we should note that it’s also quite complicated in other cases we regard as promising candidates for reduction. For example, it’s not clear what exactly is involved in *being healthy* or *being alive*, but we have some ideas of the more basic natural (physical, chemical, etc.) properties involved in these things, and further scientific investigation will hopefully bring us closer to a full understanding of these properties and how they fit in

the natural world. In a similar vein, Boyd (1988) suggested that goodness might be initially viewed as a cluster of natural properties: the ones conducive to the satisfaction of human needs, tending to occur together and with a tendency to promote each other (or to be promoted by the same sorts of things). It's not obvious what exactly is involved in this, but we surely have some idea of the sorts of things these could be. If this is right, then maybe we just need to do more scientific investigation to fully understand the nature of goodness. This would mean that the realist is as entitled to view goodness as a natural property as they are to view being healthy or being alive as natural properties. This optimism about a posteriori reduction of the moral to the natural is characteristic of one prominent form of ethical naturalism: **a posteriori reductive naturalism**.

The main objection to a posteriori naturalism comes from the **Moral Twin Earth Thought Experiment** (Horgan and Timmons 1991, 1992), which is very similar to Hare's missionary and cannibals example. The idea is to imagine two worlds that are very similar: Earth and twin Earth. The only difference is that the properties that causally regulate the use of ethical terms in these two worlds are slightly different. For example, imagine that on Earth use of the term 'right action' is causally connected to whatever maximizes overall happiness, whereas on twin Earth use of the term 'right action' is causally connected to whatever is done for reasons that one could universalize. (Basically, they differ in whether they are implicitly utilitarian or Kantian about what's right action.) Nevertheless, the practical role of this term in each world is the same: thinking that an action is right tends to motivate people to do it, people are resented and punished when they knowingly don't do the right action, and so on. Now, if the causal theory of reference is correct, then the term 'right action' refers to a different property when used on Earth and twin Earth, since subtly different properties are causally linked to use of this term in the two worlds. Intuitively, however, it doesn't seem that Earthlings and twin-Earthlings are talking about different things when they talk about what actions are right. They're imagined to be in different possible worlds, so they can't talk *to* each other. But it seems that, if they were to talk to each other and disagree about whether a particular action is right, this would be a genuine disagreement rather than a case of talking past each other. After all, thinking that an action is right plays the same practical role for both of them.

This is called the *Moral* Twin Earth Thought Experiment because it's modeled on a similar Twin Earth Thought Experiment famously used by Putnam (1975b) to show that if twin Earthlings' use of the term 'water' were causally regulated by a different chemical compound from Earthlings' use of the term 'water,' we'd be inclined to think they're talking about different stuff, even if the stuff that flows in the rivers and streams is called water in both worlds. Horgan and Timmons suggest that our intuitions go the other way in the moral case, showing that the causal theory of reference cannot be used to support a posteriori reductionist forms of ethical naturalism. It's worth noting the structural similarities between this objection to naturalism and the previous objection from disagreement against relativism.

Key Point: A posteriori naturalists think ethical facts can be reduced to natural facts, not by the sort of conceptual analysis that Moore criticized but by empirical identification of ethical properties with properties that are investigated by the natural sciences.

• A PRIORI NETWORK NATURALISM

Let's move on to a different strategy for explaining how ethical facts might be natural facts. As we saw before, traditional attempts at conceptual analysis of ethical terms in natural terms proved problematic in light of the open-question argument. However, some terms seem to hang together in a network of related concepts that likewise resists traditional analysis, but this doesn't lead us to say that they refer to nonnatural properties. For example, the words 'game,' 'rules,' 'play,' 'winning,' 'losing,' etc. seem to be closely related and to form some kind of package. We may be unable to reductively analyze any one of these in terms of the others, let alone wholly in terms not in the package. Nevertheless, some have thought it is still possible via a priori reflection to map out the network in a detailed enough way that we can identify the properties referred to by individual terms of the network with whatever property uniquely satisfies a particular role in the network.

The basic idea behind the last form of metaethical naturalism we will discuss here is that ethical terms hang together in a network like this, and mapping out the network is crucial for showing the properties they refer to to be natural. If that is right, then it shouldn't be surprising that we cannot provide an *atomistic* analysis of ethical terms in natural terms—as the open-question argument assumes the naturalist would have to do—but that won't impugn naturalism as long as a **network-style analysis** is available instead.

QU9: Why does the open-question argument assume that the naturalist would have to give an *atomistic* analysis of ethical terms in natural terms?

(Warning: what follows gets considerably more complicated than most of the rest of this book, and I suspect it will be difficult for readers without a good understanding of formal logic. So it wouldn't be unreasonable to skip the rest of this section accepting only the very vague description of a priori network naturalism given above. Alternatively one might wish to skim this section and then read some of the literature on the so-called **Canberra Plan** before coming back to study this section more carefully.)

This network strategy traces back to Ramsey's discussion of how to understand degrees of beliefs and preferences in the prediction of people's behavior, but it has been more recently developed by Jackson (1998) for ethical terms. On his way of pursuing the idea, we seek to analyze a particular concept as expressing the unique

property that satisfies a network of platitudes. Platitudes are usually thought of as claims that one's usage of a term must respect in order for it to count as revealing competence with the concept expressed.

Accordingly, a network analysis can be pursued in four steps:

- 1 List the platitudes in a standardized form ("property-name form").
- 2 Form a large conjunction of the platitudes.
- 3 Replace the concepts/properties being analyzed with free variables.
- 4 Bind the variables with existential quantifiers and add a uniqueness condition.

That's a bit jargony, but the best way to understand each of these steps is just to work through a non-ethical example (after Smith, 1994): color terms. First, we put platitudes about these into a standardized form which mentions the names of properties (this is what makes it in the "property-name form") e.g.:

Color platitudes

- Objects have the property of being red if they look to have the property of being red to normal perceivers under standard conditions.
- The property of being red is more similar to the property of being orange than to the property of being yellow.
- Objects have the property of being orange if they look to have the property of being orange to normal perceivers under standard conditions.
- The property of being orange is more similar to the property of being yellow than to the property of being green.
- Objects have the property of being yellow if they look to have the property of being yellow to normal perceivers under standard conditions.
- The property of being yellow is more similar to the property of being green than to the property of being blue.
- And so on.

Second, we form these into a large conjunction:

Objects have the property of being red if they look to have the property of being red to normal perceivers under standard conditions **AND** the property of being red is more similar to the property of being orange than to the property of being yellow **AND** objects have the property of being orange if they look to have the property of being orange to normal perceivers under standard conditions **AND** the property of being orange is more similar to the property of being yellow than to the property of being green **AND** objects have the property of being yellow if they look to have the property of being yellow to normal perceivers under standard conditions **AND** the property of being yellow is more similar to the property of being green than to the property of being blue...

Third, we strip out mention of the “property-names” and replace them systematically with what logicians call free-variables (that is, variables floating freely in the formula and not linked up with any quantifier terms such as “for all x ” or “there exists a y ”):

Objects have v if they look to have v to normal perceivers under standard conditions AND v is more similar to w than to x AND objects have w if they look to have w to normal perceivers under standard conditions AND w is more similar to x than to y AND objects have x if they look to have x to normal perceivers under standard conditions AND x is more similar to y than to z ...

Since the variables aren’t linked up with any quantifiers, that sentence doesn’t really mean anything on its own. But we can make it meaningful again by doing what logicians call “binding” the free variables with existential quantifiers which look like this: $\exists$ (for “there exists”):

$\exists v \exists w \exists x \exists y \exists z \dots$ (Objects have v if they look to have v to normal perceivers under standard conditions AND v is more similar to w than to x AND objects have w if they look to have w to normal perceivers under standard conditions AND w is more similar to x than to y AND objects have x if they look to have x to normal perceivers under standard conditions AND x is more similar to y than to z ...),

and we add a uniqueness condition to insure that there is not more than one of each property:

AND if any v' satisfies this condition then $v'=v$, and if any w' satisfies this condition then $w'=w$, if any x' satisfies this condition then $x'=x$, and if any y' satisfies this condition then $y'=y$, and if any z' satisfies this condition then $z'=z$...

This big (and quite ugly) sentence claims that there are unique properties v , w , x , y , z ... that are related in a network that would satisfy the color platitudes.

The preceding few paragraphs may seem totally complicated and arcane, but here’s the payoff of these logical shenanigans: we now have a way of explaining what redness is which doesn’t use any color terms whatsoever. We’ll say that redness is the property x such that it satisfies this big conjunction (notice that the conjunction above doesn’t have any color terms in it). Similarly, we can say that orangeness is the property w which satisfies this big conjunction. And so on, for each color, thereby providing a network analysis of color properties which doesn’t itself use color concepts in the analysis. The key idea is that, although we might not be able to analyze ‘red’ by conceptually decomposing it into its atoms (on the model of ‘bachelor’ = ‘unmarried man’), we’re still in a position to use a priori reflection on our color concepts to explain in non-color terms what redness is: it’s whatever unique property plays the relevant role in this network of concepts.

At this stage, you should be wondering how this is going to apply to the ethical case. Jackson’s idea is that here too we might start by listing various platitudes, e.g., about what’s right/wrong, good/bad, virtuous/vicious, and how they’re related to each other.

Then we could pursue steps 1–4 as before. And if successful, we'd end up with an analysis of ethical properties such as goodness and rightness that explained their nature in wholly non-ethical (and hopefully appreciably naturalistic) terms, thereby providing an a priori network-style reduction of ethical properties to non-ethical properties.

In this case, however, there's a big challenge. Morality is a much more controversial topic than color, which means that it is hard to come up with very many platitudes. (Remember that "platitudes" are usually understood to be claims that one's usage of a term must respect in order for one to have revealed competence with the relevant terms.) To be sure, there are some very minimal things we might be able to say, e.g. (in the standardized property-name form):

Moral platitudes

- If an action has the property of being wrong, then doing it does not have the property of being right.
- Encouraging others to do actions with the property of being right itself has the property of being right.
- A person with the property of being virtuous does actions with the property of being good.
- And so on.

But as Zangwill (2000) argued, it's difficult to get much more substantive than that. But unless we do, there won't be enough platitudes to have any hope of justifying the idea that some unique set of properties satisfies them.

Moreover, even if we do admit more ethical claims as platitudes, there's a risk that the a priori network analysis will still be unstructured enough to face what Smith (1994, 54–6) calls the **permutation problem**. This is the problem that, once we strip out all of the ethical terms from these platitudes and put in free variables, we could very well have a structure that can be inverted or systematically shifted (e.g., switching all positive terms such as 'good' with negative terms such as 'bad'). To see what he has in mind, imagine a "shifted" spectrum of color terms, where 'green' is used to refer to blueness, 'blue' is used to refer to redness, 'red' is used to refer to greenness, and so on. These color terms would equally well satisfy the network structure extracted from color platitudes above, but they'd seem to not provide the correct analysis of the meanings of 'red,' 'green,' and 'blue.' Similarly, we might imagine an inverted morality that satisfies the network structure extracted from the minimal list of moral platitudes above. Again this would mean that there needn't be a *unique* set of properties related in the way the network articulates. In technical terms, the "uniqueness condition" mentioned above would be unjustified.

Jackson himself tries to avoid these kinds of problems by suggesting that what should be fed into the network analysis as "platitudes" in the case of ethical properties is not merely the ethical claims that *everyone* could agree to but rather the fixed points of what he calls "mature folk morality." By this, he means something like the principles

everyone would agree to if morality was perfected via critical reflection. His hope is that this will provide a more substantive list of “platitudes” on the back of which to pursue a network analysis. Indeed, we might even end up in a similar place to Boyd’s brand of metaethical naturalism via a much different route: saying, e.g., that goodness is a cluster of natural properties conducive to the satisfaction of human needs, tending to occur together and with a tendency to promote each other (or to be promoted by the same sorts of things). If that is indeed the unique property satisfying the goodness-node in a network of concepts extracted from the fixed points of mature folk morality, then perhaps we should recognize goodness as that (complex) property.

However, one might reasonably worry that analyzing goodness in this way undermines the aspirations of a priori analysis to provide an exposition of the contents of our ordinary ethical concepts. For the platitudes were supposed to capture the commitments, the denial of which would reveal some kind of conceptual/semantic incompetence with the relevant terms. The problem is that failing at the critical reflection needed to arrive at “mature folk morality” doesn’t reveal that one is incompetent with the terms one uses in our everyday moral discourse. Moreover, as Yablo (2000) points out, appealing to *mature* folk morality may smuggle in an ethical concept into an analysis that was supposed to be entirely reductive in the sense that it explains the nature of ethical properties without appealing to ethical concepts. If “mature folk morality” is just a way of saying the morally right morality for us to appeal to, then obviously we haven’t reductively analyzed ethical concepts via the network analysis.

Key Point: A “network analysis” promises to provide an a priori reduction of ethical properties to natural properties by analyzing each ethical property as qualities that play a particular role in a network of related qualities that can be understood a priori.

• CONCLUSION

In this chapter we have explored four versions of ethical naturalism, the doctrine that there are ethical facts and they are natural. One of the complicating factors in this debate is that no one has an uncontroversial definition of what it takes for a fact to count as natural. But assuming the facts in principle discoverable by science or figuring in the law-like explanations characteristic of scientific explanation are natural facts, the challenge for the naturalist is to show how ethical facts might reduce to or fit among such facts.

The attractions in attempting to do this and the main general arguments in favor of naturalism can be put in terms of (i) metaphysics and epistemology, and (ii) philosophies of language and mind. Metaphysically, naturalism is a form of *realism*: ethical facts really obtain. But unlike the nonnaturalist, the naturalist thinks these facts fit within

the natural world. If the view can otherwise be made to work, this means that it doesn't require adding a stock of *sui generis* and poorly understood "nonnatural" facts to our overall ontology. Moreover, this holds out some promise of explaining how we come to know ethical facts without appealing to some special faculty by which we glom onto to *sui generis* facts. In the philosophy of language, naturalism is a form of *representationalism*: ethical statements are like other statements of fact in that they represent a way reality could be, and they are true or false depending on whether reality is the way they represent it as being. This is because they express beliefs about how reality is, which is why the naturalist is, in the philosophy of mind, a *cognitivist* about ethical thought. Here too, if the view can otherwise be made to work, this is attractive theoretical territory to occupy. For linguistically the sentences with which we make ethical statements seem to behave like other statements commonly regarded as representations of reality (e.g., they're truth-apt, embeddable under 'if' and 'might'). Moreover, it's quite unremarkable for us to talk about ethical beliefs and knowledge, which inclines many metaethicists towards some form of noncognitivism.

As we have seen, there are several strategies for locating ethical facts among the natural facts, and each strategy faces its own specific objections. But we might say that the main liability of naturalism lies in the philosophy of language and the philosophy of mind. Each form of naturalism here trades on controversial assumptions in the philosophy of language. Moreover, for most natural facts, such as facts discovered by chemistry or physics, stating that one of them obtains does not (all by itself) reveal in the speaker any particular motivations to act, and the obtaining of one of these facts does not (all by itself) generate reasons for someone to act. So when it comes to the nature of ethical facts, on any particular naturalist view, it's worth asking: if ethical facts were like that, would belief in them plausibly generate (defeasible) motivations? And: would a fact like that be able to generate reasons to act all by itself?

• CHAPTER SUMMARY

- Neo-Aristotelian naturalists expand the conception of the natural standardly assumed in the metaethical debate and attempt to derive ethical facts from human nature.
- It is challenging for the neo-Aristotelian naturalist to distinguish how humans biologically should be from how humans ethically should be without smuggling in ethical considerations into the determination of what counts as human nature.
- Ethical relativists challenge the aspirations of other forms of ethical realism to secure universal ethical facts. If we instead accept that ethical facts are relativized to cultures or individual value systems, it is much easier to locate such facts among the facts we already regard as natural.
- Traditional arguments for nonnaturalism assume that for ethical fact to be shown to be natural, we would have to be able to decompose ethical concepts in purely naturalistic concepts. By appealing to the difference between metaphysical identity and conceptual analyzability (e.g., water = H₂O), a posteriori naturalists point

out that this is a bad assumption and argue that there is reason for optimism that ethical properties are identical to natural properties.

- A posteriori naturalists are threatened by the moral twin Earth objection.
- Network analyses offer an alternative model for a priori conceptual analysis of ethical terms than that assumed by Moore and his critics.

• STUDY QUESTIONS

- 1 What does it take for a fact to count as “natural”?
- 2 What would count as a “virtue” of a car, a cactus, and a tiger on the neo-Aristotelian approach? Explain your answer.
- 3 Are there unethical aspects of human nature?
- 4 Are ethical relativists realists or antirealists about ethics?
- 5 Explain the causal theory of reference.
- 6 Why is the possibility of a “moral twin Earth” a threat to a posteriori naturalists?
- 7 What is an example of a moral platitude; would platitudes like this support a network analysis of ethical terms?

• FURTHER RESOURCES

- Gowans, Chris. 2015. “Moral Relativism,” *The Stanford Encyclopedia of Philosophy* (Summer 2015 Edition), Edward N. Zalta (ed.), <http://plato.stanford.edu/archives/sum2015/entries/moral-relativism>. [A detailed general introduction to various forms of relativism in ethics.]
- Harman, Gilbert and Thomson, J. J. 1996. *Moral Relativism and Moral Objectivity*. Blackwell. [Book-length debate between a relativist and an objectivist about morality.]
- Hursthouse, Roslind. 1999. *On Virtue Ethics*. Oxford University Press. [A book-length development of the neo-Aristotelian approach to ethical naturalism.]
- Lenman, James. 2014. “Moral Naturalism,” *The Stanford Encyclopedia of Philosophy* (Spring 2014 Edition), Edward N. Zalta (ed.), <http://plato.stanford.edu/archives/spr2014/entries/naturalism-moral>. [A detailed general overview of naturalist positions in metaethics outside of relativism.]
- Papineau, David. 2015. “Naturalism,” *The Stanford Encyclopedia of Philosophy* (Fall 2015 Edition), Edward N. Zalta (ed.), <http://plato.stanford.edu/archives/fall2015/entries/naturalism/>. [Section 2.3 contains further discussion of the “Canberra Plan.”]

• ANSWERS TO QUESTIONS OF UNDERSTANDING

QU1: (i) and (iv) are attributive.

QU2: Cars and lamps are functional kinds, moons are not (because it’s not part of being a moon that it has any particular function). The case of human beings

is the interesting one—do we have a particular function (or purpose?) which is essential to being human? Some have argued yes and others have argued no. Crucially for what follows, you don't have to think humans have a particular function to pursue neo-Aristotelian naturalism, but you do have to think that humans have something from which we can derive what it is to be a good human.

- QU3: Your answers to this question may differ, as there are often many different features that make something which is a functional kind a good instance of its kind. For example, shoes are good shoes when they're comfortable, durable, fashionable, etc. Power drills are good power drills when they have sufficient torque, are easily gripped, have variable speeds, etc. And, perhaps, textbooks are good textbooks when they have questions of understanding such as this one.
- QU4: First, try to come up with an account of the characteristic way of living for a dairy cow. Next, figure out which features of a dairy cow facilitate its excelling at this characteristic way of living. These are the features a dairy cow (in light of it being a dairy cow) should have.
- QU5: There may be more than one in each sentence, but at least the following: (i) is relative to something like a degree of precision or a practical purpose, (ii) is relative to a kind of trip, and (iii) is relative to what type of person we're comparing Sarah to.
- QU6: To understand an utterance of sentence (i) we need to know *which direction* determines front and back. To understand an utterance of sentence (ii) we need to know when counts as "today" relative to which "yesterday" is the previous day; we also need to know where is "here" for the utterance.
- QU7: However carefully we examine our concepts, we won't discover that what it is to be water is to be H₂O; for that we need to do empirical investigation of samples of water, applying the tools of modern chemistry.
- QU8: The name 'Gödel' still refers to the man named Gödel. It's just that he turned out not to have been the mathematician who proved what we now call "Gödel's incompleteness theorem."
- QU9: Because the open-question argument demands an analysis of a term such as 'good' that breaks it down into other terms which are natural. This is why Moore is sometimes accused of confusing the nonnaturalness of goodness with the simplicity of goodness.

• WORKS CITED

- Boyd, Richard. 1988. "How to Be a Moral Realist." In *Essays on Moral Realism*, edited by G. Sayre-McCord. Ithaca, NY: Cornell University Press.
- Dreier, James. 1990. "Internalism and Speaker Relativism." *Ethics* 101 (1): 6–26.
- Foot, Philippa. 1999. *Natural Goodness*. Oxford: Clarendon Press.
- Geach, P. T. 1956. "Good and Evil." *Analysis* 17 (2): 33–42.
- Harman, Gilbert. 1975. "Moral Relativism Defended." *The Philosophical Review* 84 (1): 3–22.

- Horgan, Terence, and Mark Timmons. 1991. "New Wave Moral Realism Meets Moral Twin Earth." *Journal of Philosophical Research* 16: 447–65.
- . 1992. "Troubles on Moral Twin Earth: Moral Queerness Revived." *Synthese* 92: 221–60.
- Hursthouse, Roslind. 1999. *On Virtue Ethics*. Oxford: Clarendon Press.
- Jackson, Frank. 1998. *From Metaphysics to Ethics*. Oxford: Oxford University Press.
- Kripke, Saul A. 1980. *Naming and Necessity*. Cambridge, MA: Harvard University Press.
- Putnam, Hilary. 1975a. *Mind, Language, and Reality*. Cambridge; New York: Cambridge University Press.
- . 1975b. "The Meaning of 'Meaning'." In *Philosophical Papers, Vol. II: Mind, Language, and Reality*. Cambridge: Cambridge University Press.
- Smith, Michael. 1994. "Internal Reasons." *Philosophy and Phenomenological Research* 55 (1): 109–31.
- Yablo, Stephen. 2000. "Red, Bitter, Best." *Philosophical Books* 41: 13–23.
- Zangwill, Nick. 2000. "Against Analytic Moral Functionalism." *Ratio* 13 (3): 275–86.

• NOTES

- 1 With the possible exception of some forms of quasi-realist expressivism. When this view was first developed by Blackburn in the 1980s and early 90s, it appealed to minimalism about 'true' and 'believes' to earn the right to use some of the realist sounding language about ethics that is part of ordinary everyday ethical discourse. One way the view has been developed, however, lets this minimalism extend to any metaphysical terminology, whether or not it is part of ordinary everyday ethical discourse. If this also includes the idea of "reality including facts" of some type, then strictly speaking, the quasi-realist needn't accept (C) as it is stated above. I'll ignore this complication in what follows.
- 2 It's a controversial issue in metaphysics how exactly to understand such "natures." The concept is related to but more narrow than the idea of an essence.
- 3 Sometimes philosophers distinguish between a *descriptive* claim about morality (roughly, its demands vary across cultures) and a *normative* ethical claim (we should respect cultural differences and not expect that all people live by the same basic moral rules). It's important to note that the form of relativism discussed above is the descriptive sort. The normative claim is ostensibly a *universal* ethical claim—so if it is true, then it undermines the descriptive relativist's claim that there is no universal moral truth.
- 4 In Chapter 7, we'll encounter response-dependence and constructivist theories that hold that there are ethical facts but insist that they are not fully "out there" because their ontological status depends somehow on the responses or constructions, actual or hypothetical, of some kind of mind. It's a controversial issue in metaethics whether to view these as a sophisticated form of realism, not quite conforming to the standard definitions, or to view these as sophisticated forms of antirealism. I'll treat them as antirealist proposals in Chapter 7, but I do so in the context of a discussion about whether some of the standard categories begin to break down in various ways in more recent metaethical debate.