

8

• outstanding issues

At the beginning of this book, we considered ethical judgments such as the one made with the statement “Bush’s invading Iraq was wrong.” We haven’t ventured any opinion here on whether such judgments are correct. Instead we’ve investigated the status of these judgments. This has involved exploring issues in metaphysics, epistemology, philosophy of language, and philosophy of mind as they apply to ethical thought and discourse. That is, we’ve focussed on metaethics.

One of the exciting recent developments around metaethics is the extension of the discussion of metaphysical, epistemological, linguistic, and psychological issues to other domains. For example, much like one can say that Bush shouldn’t have started the war in Iraq, one might say that he shouldn’t have believed that there were caches of weapons of mass destruction. However, evaluating beliefs in this way is not (normally) ethical; rather it’s “epistemological.” That is, it’s about whether the belief meets *epistemic* standards for good belief, where these are usually based on evidence and truth rather than morality.

In a similar vein, we might ask whether someone’s belief is epistemically right/wrong or good/bad, and also about whether a believer is epistemically virtuous/vicious. This is not the same as, but it is analogous to, how we ask about whether actions are ethically right/wrong or good/bad, and also about whether agents are ethically virtuous/vicious.

Thought and discourse about such matters is epistemological thought and discourse. And it would seem that all of the same metaphysical, epistemological, linguistic, and psychological questions we have been exploring about ethical thought and discourse could also be raised about such *epistemological* thought and discourse. In doing so, we are no longer doing metaethics; we’re doing **metaepistemology**.

Key Point: Metaepistemology is to epistemology and epistemological thought and discourse as metaethics is to ethics and ethical thought and discourse.

Broadening our focus in this fashion suggests a further interesting question. What do ethics and epistemology have in common such that they're species of some genus; and are there any other species of this genus? The short answer we've already encountered in previous chapters is that the genus is *normativity*. Just like we might try to distinguish normative ethics (first-order) from metaethics (second-order), we might try to distinguish "normative epistemology" from metaepistemology. But there seem to be other normative domains. For example, focussing only on what would best promote a particular person's well-being, we can ask about what actions that person *prudentially* ought to perform. Prudential normativity is arguably different from (though connected to) ethical and epistemological normativity.

We can also ask about what actions would be rational to pursue given some stipulated end given the agent's background beliefs. That is, we can ask about what is **instrumentally** right/wrong, good/bad, etc. for achieving a stipulated end (which needn't be ethical or prudential). So, instrumentality (perhaps as a kind of rationality) might be viewed as another domain of normativity beside ethics, epistemology, and prudence. Moving away from actions and beliefs, we can also ask about what emotions or feelings are reasonable, rational, or justified to have in various circumstances. Sometimes this is an ethical question, but it's not always a matter for ethics. Should I feel elated if my team wins the championship? Probably so, but that's not really an ethical matter, yet it is an issue about a "should" and so the "fittingness" of attitudes might be conceived as another species of normativity. In a somewhat similar vein, focussing only on considerations about beauty and style, we might consider what is aesthetically right/wrong, good/bad, etc. So aesthetics too might be viewed as another domain of normativity.

With each of these domains of normativity, we can imagine a meta-level investigation into their status, i.e. an analog to metaethical investigation but into the metaphysics, epistemology, philosophy of language, and philosophy of mind for these other normative domains. Moreover, this suggests a further development in recent metaethical discussions: what happens to the debate when we broaden the focus to normativity *in general*? As before, rather than asking first-order normative questions, this is conceived as a second-order investigation into questions about metaphysics, epistemology, philosophy of language, and philosophy of mind as they apply to our first-order normative thought and discourse, in its full generality. It is controversial how exactly to use the technical term "normativity," but whatever the various domains of normativity are, this represents an expansion of metaethics into what is sometimes called **metanormative theory**.

QU1: What are some terms or concepts that apply to multiple normative domains?

In this chapter, we shall start with metanormative theory and consider the question of how those favoring one of the four traditional metaethical theories might react to

the expanded focus on normativity in general. What we'll see is that many (though not all) of the arguments for and against those traditional theories can be refashioned as arguments for and against similar theories about so-called **all-things-considered normativity**. This will provide a review of many of the arguments of Chapters 2–5. Then we'll consider one specific domain of normativity that is not ethics: epistemology. We'll explore some of the theoretical options and associated costs and benefits in order to give you a sense of how a similar debate might take place with respect to a specific non-ethical but still normative domain.

As these are areas of philosophical inquiry that are only now beginning to receive much attention, this chapter is intended as a very open-ended conclusion that will hopefully put you in a position to see how many of the issues we've covered in previous chapters might be extended and adapted to other domains in exciting and novel ways. Just as metaethics largely involves a fruitful application of other areas of philosophy to the ethical domain, we'll begin to see here how metaethics can be extended and morphed in ways that benefit our understanding in other areas of philosophy.

• FROM METAETHICS TO METANORMATIVE THEORY

Consider again the statement “Manning shouldn't have released classified documents to Wikileaks.” We've investigated various metaethical positions one might take about the status of this statement considered as the expression of an *ethical* judgment. Nonnaturalism, expressivism, error theory/fictionalism, naturalism all constitute families of views about the status of judgments like this one. There is, however, a way to hear this statement where it isn't specifically ethical. For instance, if someone points out how bad it has turned out for Manning (e.g., how she will now spend many years in prison), this person may be neutral on the *ethics* of releasing the classified documents but still insist that it was a *prudentially* very bad thing to do. That is, someone might think that, whether or not it was ethically right, releasing the documents was the wrong way for Manning to promote her own well-being.

This distinction raises an interesting and difficult issue. Granting that ethics and prudence aren't the same thing, one might now wonder what Manning should have done *all things considered*. That is, considering every relevant factor from ethics to prudence to anything else that might be relevant, what should she have done?

More generally, this question is sometimes put by asking what someone *just plain* ought to do. Ethics and prudence concern norms (paradigmatically) for action, but it is far from clear how they interact with each other and with other norms for action. Some philosophers think ethics trumps other norms. So, on this view, a prudential reason could add to an ethical reason to give someone even more reason to do something, but if ethics and prudence point in opposite directions ethics always wins. Other philosophers, however, hold the opposite view: prudence trumps ethics.

A third possible position is that there is no trumping normative domain: ethics and prudence (and perhaps other kinds of normativity) contribute reasons bearing on action, and it requires careful practical reason to figure out what we “just plain” ought to do. Determining the right view here is an interesting and difficult question at the interface of normative and metanormative theorizing.

Given our focus earlier in this book on issues in metaphysics, epistemology, philosophy of language, and philosophy of mind, metanormative theorists will also want to ask how these meta-level issues extend to all-things-considered normativity. As this is such a new area of philosophy, there are not that many positions to consider. But one thing we can do is to ask how each of our traditional metaethical theories would extend to all-things-considered normativity.

Nonnaturalism

Let’s begin as we did before with nonnaturalism. One of the central reasons some philosophers find this position attractive in metaethics is that they think our ethical thought and discourse seems to be fact stating and answerable to something objective but its role in our deliberations seems to be fundamentally different from the role of thought and discourse about the natural (or supernatural) world. So these philosophers draw the conclusion that ethical facts must be *sui generis*.

As I pointed out in Chapter 2, by “*sui generis*” nonnaturalists should not be interpreted as meaning that ethical facts are fundamentally different from other *normative* facts. Now we’re in a position to see why. The same motivation one might have for thinking that *ethical* facts are nonnatural can lead one to think that *normative* facts are nonnatural. When deliberating about what to do, we reach conclusions that can seem to be answerable to something objective but not reducible to the facts of the natural (or supernatural) world. If a nonnaturalist endorses the view that ethical norms are trumping, then nonnaturalism about all-things-considered normativity is almost automatic. But even if one endorses one of the other views about the relation between ethical normativity and all-things-considered normativity, the basic motivation towards nonnaturalism remains the same. And most contemporary metaethical nonnaturalists are also *metanormative* nonnaturalists.

QU2: When nonnaturalists say that ethical and normative facts are *sui generis*, with what other kinds of facts are they contrasting ethical and normative facts?

Notice, moreover, that the three more specific arguments we considered for nonnaturalism about ethical facts seem not to trade on anything specifically ethical but rather work equally well when refashioned to be about all-things-considered normativity. The first argument came from Hume’s Law, which recall is the idea that you

cannot derive an 'ought' from an 'is.' If this is true, it would seem to apply just as much to the all-things-considered 'ought' as it does to the ethical 'ought.'

The second argument was Moore's open-question argument, which traded on the apparent semantic gap between ethical words such as 'good' and 'right' and any proposed analysis in purely naturalistic terms. Interestingly, if we restrict consideration to what is good for promoting some particular end, or what action is the instrumentally right action given one's evidence, then it's easier to imagine that facts about this are natural facts. They're just facts about probabilities, arguably discoverable by empirical science. However, when applied to all-things-considered normativity, the open-question argument would seem to work as well or even better as it does in ethics. After all, one of the tricky issues about all-things-considered normativity is how exactly each different species of normativity ought to be balanced. So there is even more scope for a semantic gap between naturalistic terms and normative terms.

Finally, the last argument we considered was the deliberative indispensability argument. Roughly, this was the idea that we have to assume the existence of facts about what we ought to do in order to make sense of the way we deliberate practically. In Chapter 2, we saw how this argument actually doesn't directly establish nonnaturalism about *ethical* facts, as it is not obvious that we have to assume the existence of specifically ethical norms to make sense of practical deliberation. If anything, it establishes nonnaturalism about all-things-considered normative facts (and extends to ethical facts because of their centrality to general normativity).

Just as much as the motivations towards a nonnaturalist view in metaethics extend to metanormative theorizing, the arguments against it remain powerful. We considered four main objections: the objection from a naturalistic worldview, epistemological objections, challenges from the practicality of ethics, and an argument from supervenience.

Surely, if one is inclined to reject nonnaturalism about ethics because of one's commitment to a naturalistic worldview, expanding our focus to all-things-considered normativity is unlikely to make one any less skeptical of adding new *sui generis* kinds of facts to our ontology. As I explained in Chapter 2, a more specific form of this objection to nonnaturalism derives from considerations in epistemology about how we could come to know nonnatural facts; and this applies to all-things-considered normative facts as much as it does to specifically ethical facts. If these are not knowable via a posteriori observation and a priori reflection, we will have to posit some new faculty by which we know them. This imposes a heavy explanatory burden on the nonnaturalist to work out a plausible view of where this faculty comes from and how it works. Moreover, many philosophers will think that judgments about what I "just plain" ought to do are surely as action-guiding if not more than judgments about what I *ethically* ought to do. That's not to say that any of these three objections provide "knockdown arguments," but they do seem to extend from ethics to all-things-considered normativity.

The supervenience objection to nonnaturalism is more delicate to extend from meta-ethics to metanormative theory. Recall that the term supervenience refers to the idea that issues in one domain (e.g., how much punishment some kind of criminal deserves) cannot vary without some underlying variation in another domain (e.g., the facts about what they did). Most philosophers think ethical normativity supervenes on natural fact, because otherwise the ethics of a situation would seem arbitrary. For similar reasons, many philosophers will also think what one just plain ought to do also supervenes on natural fact. There may, however, be a kind of existentialist viewpoint according to which at least some question of what one ought all-things-considered to do are systematically thwarted by a kind of radical freedom, forcing unreasoned choice rather than norm following. The very intelligibility of this viewpoint makes the supervenience argument a bit more difficult to apply to all-things-considered normativity. However, notice that most nonnaturalists wouldn't want to grant that sometimes what one just plain ought to do is a matter of radical choice rather than belief tracking some external fact—for then they lose some of the motivation for thinking that there are such facts. Moreover, this move would seem to concede something important to the expressivist who thinks such judgments are more like intentions or plans than like beliefs about some piece of reality. Anyway, it is not obvious that the existentialist viewpoint mentioned here is intelligible in the end, so these issues remained vexed.

Expressivism

Expressivism, you'll recall, began its life in the form of emotivism, which said that ethical statements are the expression of emotive reactions to reality rather than representations of reality. This may seem more plausible for ethical normativity than for normativity in general. After all, it often does seem that our ethical opinions are laden with emotive affect, but it is less obvious that one's opinion about what to do all things considered is similarly emotive.

It's important to remember, however, that emotivism was only one early version of expressivism, and the expressivist position became much more sophisticated as it developed over the twentieth century. And contemporary expressivists tend to see their view as one about all-things-considered normativity which has applications in domains such as ethics and epistemology. Keeping that in mind, let's now consider whether the main arguments given for expressivism in the ethical case extend to the case of normativity in general.

The first argument we explored before was basically a twist on the argument from Hume's Law for nonnaturalism. The idea was that if one thinks Hume's Law provides a reason to doubt that ethical facts can be reduced to natural facts, but one is deeply committed to a naturalistic worldview, then one won't want to endorse nonnaturalism, and expressivism seems like the main alternative. For expressivists can respect a fundamental difference between "ought" and "is." They chalk it up not to two different

kinds of facts but rather to two different roles statements can play in our language: descriptive statements about how reality is and emotive/prescriptive statements about what people ought to do. This argument has nothing specifically to do with ethics and so works equally well when applied to all-things-considered normativity.

The second argument for expressivism we considered in Chapter 3 was the argument from motivational internalism and the Humean theory of motivation. The idea was that there is a specially tight or “internal” connection between ethical judgments and motivation to action. And assuming motivation to action always requires the cooperation of two kinds of mental states, one representational and the other directive, the most plausible explanation of the internal connection between ethical judgment and motivation is that these judgments are more desire-like (directive) than belief-like (representational). As we discussed in Chapter 3, the internalist thesis about ethical judgments is very controversial. Some have argued that there could be psychopaths or amorlists who make ethical judgments in full awareness of what they mean but who lack all motivation to act in their accord. If anything, this objection is less forceful when we turn to all-things-considered normativity. For it is more difficult to imagine the analog of a psychopath or an amorlist about what one “just plain” ought to do. However, another worry that we raised for this argument was that we often make ethical judgments about things very different from us or far removed from us, and it is far less clear that these bear the same “internal” connection to motivation as ethical judgments about our own prospective courses of action. This worry applies to a similar argument from internalism plus the Humean theory of motivation for expressivism about all-things-considered normativity.

The final main argument we considered for expressivism in the ethical case was the argument from supervenience. The basic idea was that ethical right/wrong, good/bad, virtuous/vicious appear to supervene on non-ethical facts, and expressivism provides a strong explanation of this appearance. Insofar as supervenience counts as a theoretical cost for the metaethical nonnaturalist, it is a theoretical benefit for the metaethical expressivist. Many philosophers will think the situation is exactly the same as regards all-things-considered normativity. However, as I mentioned above, some may think that the intelligibility of a radical existentialist viewpoint about what one all-things-considered ought to do reveals a difference for this issue when extended to all-things-considered normativity. If that’s right (and this is a big “if”), then the supervenience argument for expressivism wouldn’t be as powerful in the case of all-things-considered normativity. On the other hand, as I also mentioned above, if this viewpoint is intelligible, that might play into the expressivists’ hands, since it would provide some support for the idea that at least some normative judgments are more decision-like than belief-like.

QU3: What would it be for all-things-considered “ought” to supervene on the nonnormative?

The arguments against expressivism in the ethical case extend to the case of all-things-considered normativity. In Chapter 2, I presented these and explained how they have motivated adaptations (improvements?) in the expressivist position since its early emotivist form. For example, emotivism seems to many philosophers to face difficulties making sense of ethical disagreement and the way we take ethics seriously. But prescriptivist, projectivist, and norm-expressivist versions of the view do better on this accord. These seem to be just as powerful worries in the case of all-things-considered normativity as they are in the ethical case. So expressivists wanting to extend their view will want to use one of these more sophisticated forms of the view. Same too regarding the truth-aptness (and associated features) of normative language and the possibility of embedding normative sentences in conditionals, interrogatives, etc. Here, the quasi-realist program instigated by Blackburn has been to explain how an expressivist can make sense of these features of normative discourse. Blackburn and those who have followed him in developing this program have always been concerned with more than just ethical normativity, seeking to extend their view to all forms of normativity. Whether they can successfully do so remains one of the key challenges to expressivism about all-things-considered normativity.

Error theory/fictionalism

Let's turn then to error theory and fictionalism. Mackie's error theory, which we discussed in Chapter 4, was the view that ethical thought and discourse is based on an error and should therefore be eliminated. Ethics presupposes the existence of intrinsically motivating and objectively prescriptive values, but there are no such things—or so Mackie thought, which is why he argued that all basic ethical statements are false. As we discovered, part of his reason for thinking this seems to have been based on the assumption that reasons for people to do things must be connected to those people's particular desires, cares, and concerns. But Mackie thought that ethical facts would have to be facts that generated reasons for action independently of people's particular desires, cares, and concerns.

This may suggest that there is little prospect for extending the error theory to all-things-considered normativity. For part of Mackie's argument for the error theory about ethics is based on an assumption about what reasons for action are like. So, even if he denies the existence of *ethical* reasons, he must think that there are *some* kind reasons; and so he must think that some all-things-considered normative ought-claims are true. Moreover, Mackie drew an eliminativist lesson about what we should do in light of his view that ethics is based on an ontological mistake: we *should* stop talking about ethical right/wrong, good/bad, etc. and shift our practice to talk about the sorts of reasons that there actually are for acting in various ways (i.e. the considerations attaching to people's particular desires, cares, and concerns). But surely he must think this should-claim is true, and plausibly it is intended as an all-things-considered normative claim. So even if the error theory is correct in the ethical domain, it doesn't appear to be extendable to the more general normative domain.

That may ultimately be right, but the line of thought just rehearsed can be resisted. For one thing, although Mackie's argument for the error theory drew on a view about the nature of practical reasons, not all arguments need to do this. One might think, for example, that any claim about there being reasons for someone to do something presuppose free will and then deny the existence of free will. This would imply that no claim about there being reasons for someone to do something could be true. That wasn't Mackie's argument, but it is an argument for error theory which would allow the view to be extended to all-things-considered normativity.

Moreover, the practical lesson Mackie recommends in light of his error theory about ethics is not the only lesson one could draw from the idea that all basic ethical statements are false. As we saw in Chapter 4, some philosophers embrace a form of fictionalism (speech-act fictionalism) which agrees with Mackie that all basic ethical sentences, if they were used literally to make assertions, would be false; but these philosophers argue for a more generous interpretation of ordinary ethical discourse than the interpretation Mackie embraced. These fictionalists interpret normal use of an ethical sentence as performing some other speech-act than assertion: e.g., pretense, moralizing, speaking-relative-to-a-fiction. Alternatively, some error-theorists have recently argued that even though basic ethical sentences are used to make assertions and these are literally false, there is good reason to retain the practice of making ethical assertions (Olson, 2010, Chapter 9). After all, we often have good reason to say false things. So we should conserve ethical thought and discourse in spite of its systematic basis in ontological error. Either way, if we adopt a fictionalist or **conservationalist** stance to the putative error in ethical thought and discourse, then we might extend this stance to all-things-considered normativity. For we might think that even the non-ethical "should" involved in claims about what we should do if the error theory is correct could be interpreted without contradiction along fictionalist or conservationalist lines.

QU4: What would speech-act fictionalists say about the claim that we ought (all-things-considered) to continue using all-things-considered normative discourse?

So it does seem possible to embrace the error theory not only about ethics but about normativity more generally. Of course, for anyone who objected to the error theory about ethics, this may simply amplify their worries. The main worry we discussed in Chapter 4 was based on the principle of charity. This was the idea that, when interpreting other people, we should assume that most of what they say is true; and when our interpretation violates this assumption, the burden of proof is on us to show that they really mean what we say they mean by what they say. So if it's not just lots of ordinary ethical statements that are claimed to be false but rather lots of more general normative statements, then there will be an even stronger burden to reconsider our interpretation of ordinary discourse. As far as that goes, however, it may just push us more in the direction of some kind of fictionalism about all-things-considered normativity.

We discussed two different forms of fictionalism in Chapter 4: meaning fictionalism and speech-act fictionalism. The meaning fictionalist thinks that moral statements are genuine assertions but their content is implicitly relativized to some fiction. The speech-act fictionalist, as we've just seen, thinks that moral statements aren't really assertions, but rather some other kind of speech-act. These are both positions that someone initially inclined towards error theory might take about all-things-considered normativity.

It's worth noting, however, that the main worries we considered about each theory seem to get stronger when the theories are extended. Meaning fictionalism requires thinking that ordinary users of the target discourse are "semantically blind," which means that they don't really understand the content of what they are saying. It's not that they could not articulate it explicitly (we don't expect ordinary speakers to be linguists in their ability to do semantics), but rather that they wouldn't even recognize the sort of implicit relativity to a fiction in their linguistic behavior that the fictionalist posits. So any attempt to extend meaning fictionalism from the ethical domain would require positing even more semantic blindness.

The objection about semantic blindness is part of the reason, I suggested in Chapter 4, that speech-act fictionalism is more prominent in contemporary meta-ethics. Moreover, even if we don't ordinarily talk about ethical statements as pretense or speaking-within-the-fiction, it is not uncommon to think that ethical thought and discourse is serving a "moralizing" rather than descriptive purpose. However, the challenge to identify the precise speech-act becomes all the more important and difficult as soon as the fictionalist claims that all-things-considered normative discourse involves a different kind of speech-act from assertion.

Naturalism

We discussed several different kinds of naturalism in Chapter 5. I won't review all of them here or discuss how they might be extended to all-things-considered normativity. Suffice it to say that each kind of naturalist about ethical normativity will probably want to extend their view to all-things-considered normativity, but the prospects for doing so depend on the particulars of each argument for and against each of the versions of naturalism. To give you an initial impression of this terrain, it may prove useful to consider just two prominent ways one might get to a naturalist view about all-things-considered normativity.

The first is often described as Humean because of how Hume thought reason is the slave of the passions, and the way this has been interpreted along the lines of internalism about justifying reasons as meaning that there are no reasons for action independent of agents' particular desires, cares, and concerns. (You should note, however, that this interpretation of Hume is controversial, and other metanormative views might also claim Hume as an inspiration.) If you think this, and you think that

having a desire, care, or concern is a perfectly natural phenomenon of human psychology, then you might argue that reasons to do things are simply the consequence of perfectly natural facts about human psychology.

For example, if I want to eat some ice-cream, and my freezer contains a carton of ice-cream, then plausibly I have a reason to open my freezer. A naturalist might say that this reason simply is the fact about my desire combined with the fact about where the ice-cream is. Those are perfectly natural facts, and so facts about reasons of this kind can be perfectly natural. So far, of course, that is not *all-things-considered* normativity, since I might have other reasons against opening my freezer. For example, if I want to stick to my diet and opening my freezer will tempt me to break my diet, then plausibly I have a reason *not* to open my freezer. Of course, a naturalist will insist that the fact that I have this reason is also a perfectly natural fact. To get to a Humean naturalist view about all-things-considered normativity, however, we also need a story about how reasons such as these (and *any* others that there may be) combine to determine what one just plain ought to do. We can get there with two controversial claims. First, these kinds of reasons—those made out of facts about what someone wants and facts about what has to happen in the world to satisfy this desire—are the only kinds of reasons there are for doing something. So, if these “Humean” kinds of reasons are natural, then all reasons are natural. Second, there is a plausible naturalistic account of the weighting of various reasons to determine what someone ought all-things-considered to do. Since desires come in different strengths and there are different degrees of probability that some action will satisfy some desire, perhaps a naturalist can combine these into an account of how reasons for and against any action weigh together to determine for each situation what someone just plain ought to do.

Like I said, these claims are controversial and developing the Humean naturalist view about all-things-considered normativity will require theoretical work. However, in some ways, the ethical case provides the biggest challenge to this project, as ethical norms are precisely the ones that are most often thought to provide reasons extending beyond people’s particular desires, cares, and concerns. For this reason, the subjectivist form of metaethical relativism that we discussed in Chapter 5 might seem like a good starting point for developing a Humean naturalist view about all-things-considered normativity. If ethical facts are ultimately construed as facts about what one’s own ethical values dictate, and these values count as a special class of one’s desires, cares, and concerns, then the Humean naturalist would have overcome one of the major hurdles to her view. Of course, the main worry about metaethical subjectivism was that it couldn’t make sense of moral disagreement. And this worry may seem to suggest a less radical form of relativism: one which allows for overlapping or shared values to make sense of disagreement. Still, if a naturalist can develop an account of how people’s particular desires, cares, and concerns can interact to form an overlapping or shared set of communal values, then perhaps these could still be seen as perfectly natural. Indeed, one needn’t reject the possibility of universal norms to be a Humean naturalist. For as long as one could make sense of the possibility of

there being reasons to do some things that any particular set of desires, cares, and concerns would support, one might argue that these are still perfectly natural facts about humans and the way they are situated in a world which has to cooperate for them to satisfy their desires, cares, and concerns.

The other kind of naturalist view about all-things-considered normativity I want to mention is one based not in facts about people's particular desires, cares, and concerns, but rather facts about human nature. As you'll recall, neo-Aristotelians argue that humans like many other natural systems have functions or "characteristic ways of living," which are central elements of our natures. And these philosophers seek to ground ethical facts about how we ought to act in facts about how it is good (or virtuous) for us to be, which in turn they want to ground in human nature. In Chapter 5, I presented this view as one about ethical normativity, but proponents of this view often suggest that it is a mistake to expect any sharp distinction between ethical normativity and prudential normativity. If that's right, then we might begin to think of this as a view about all-things-considered normativity. That is, we might think that what one "just plain" ought to do is grounded in one's nature. Concern for one's self, family, community, humanity, other living beings, nature, etc. might all be part of this. And the precise weighting of the reasons deriving from each of these sources might itself be something to be settled by an appeal to our natures.

That would constitute a neo-Aristotelian view about normativity in general (from which neo-Aristotelianism about ethical normativity would be an upshot). As I mentioned in Chapter 5, one of the main challenges to the neo-Aristotelian view is finding a conception of human nature meeting two constraints apparently in tension. First, this has to be a conception of human nature supporting norms that look to be genuinely ethical norms and not just biological or anthropological regularities. However, second, it must also be a conception of human nature plausibly thought to be an element of the natural world, discoverable in principle by empirical observation and scientific theorizing rather than via some *sui generis* faculty of intuition. In some ways, rejecting the sharp distinction between ethical normativity and prudential normativity might seem to help the neo-Aristotelian to address this worry. However, it remains to be seen whether their account of human nature can support intuitive claims about what we just plain ought to do. This is both because we can easily imagine situations where these point in different directions from what is biologically or anthropologically normal and because it's not obvious that there is any unified core to all of human nature.

Key Point: For each of the four main metaethical traditions, we can easily imagine parallel views in metanormative theory. The arguments for and against often look similar, but the theoretical terrain isn't exactly the same.

• FROM METAETHICS TO METAEPISTEMOLOGY

Hopefully the discussion above of normativity in general and how our four main metaethical traditions might be developed into views about what one ought, all-things-considered to do, will help you to see how metaethics is evolving and extending its reach in the recent literature. Another exciting development is the way metaethics is being used as an analogy to investigate meta-level questions about other specific normative domains. Above I mentioned epistemology, prudence, instrumentality, emotions, and aesthetics as possible further normative domains. We won't explore all of these, but I hope it will aid understanding of some of the possibilities to delve a bit deeper into metaepistemology.

Traditionally, epistemologists sought to provide an analysis of knowledge, where this was assumed to be an analysis of the *concept* of knowledge that would tell us something about the nature of the *relation* someone stands in when they know a proposition. The classic "JTB" analysis of knowledge can be viewed in this light. It breaks down the concept of knowledge into the concepts of justified, true, belief. So, for instance, if someone says that Bush knows that there are weapons of mass destruction in Iraq, then on the JTB analysis this epistemic claim is analyzed into three parts: (i) Bush believes this, (ii) it is true, and (iii) Bush's belief that it is true is justified.

But what is it for a belief to be *justified*? With this question, we begin to see an entry to metaepistemological issues. For the concept of being justified is a normative concept. And a first-order epistemological theory of justified belief is like a first order ethical theory of right action—it's a *normative* theory. This means that it's a theory about when some normative concept applies.

So, although they have rarely been viewed in this way, we can view many traditional epistemological theories as providing accounts in **normative epistemology**, viewed as a partial analog to normative ethics. For example, a rationalist theory of justification that says that one's belief must be based on good reasons in order to count as justified could be viewed as similar to a Kantian theory of right action that says that one's action must be based on good reasons in order to count as the right thing to do (though, of course, these theories may differ about when reasons are "good").

The parallel between normative ethics and normative epistemology is not perfect. But what's relevant for our purposes here is not the detail about any specific epistemological theories of justification and knowledge, but rather the fact that one can ask similar meta-level questions about the claims of these theories to the meta-level questions we asked about the claims of normative ethical theories. Not only that; even outside of the theoretical discussions of philosophy, it's a commonplace to talk about whether a belief is reasonable, justified, or rational. This is clearly similar to the way we talk about actions as reasonable, justified, or rational. Because of this, we can also ask questions about the status of ordinary epistemic thought and discourse that

are analogous to the questions we have been exploring earlier in this book about the status of ordinary ethical thought and discourse.

So, whether it is the theoretical claims of normative epistemology or the ordinary claims of everyday epistemic discourse, philosophers may wonder what we're up to when we make epistemic judgments. That is, we may wonder about the status of claims about beliefs being justified and about people knowing things. And this is the beginning of metaepistemology, much like similar questions served as the beginning of metaethics.

This is only the beginning though. As before, we can make questions about the status of some normative domain more precise by dividing them into questions of metaphysics, epistemology, philosophy of language, and philosophy of mind. In this case, that is, we can ask about the metaphysics of knowledge and justification: are these objective properties "out there" in reality; and if so, what is their nature? We can ask about the epistemology of epistemology: when we know that someone knows something, how do we know this? We can ask about the meaning of epistemic statements: are these representations of reality or expressions of attitude, or something else? Finally, we can ask about what's going on in the mind of someone who makes an epistemic judgment: does judging that I ought to believe *p* automatically motivate me to believe *p*?

Metaepistemological realism

In large part, philosophers have only recently begun to address these meta-level questions about epistemic thought and discourse. So it's much harder to outline a set of canonical metaepistemological viewpoints comprised of interlocking commitments in these four domains. However, I think it is fair to say that the traditional project of analyzing knowledge, e.g., as justified true belief, was based on an assumption that knowledge and justification are real relations somehow "out there" comprising a part of objective reality. So this is a kind of metaepistemological realism, which would go hand-in-hand with an assumption that epistemic statements are representational and epistemic judgments are belief-like rather than desire-like.

Once in the domain of realism, the immediate next question we should ask ourselves is metaphysical: do the relevant pieces of reality fit within the naturalistic worldview? In the epistemological case, this amounts to asking whether relations such as knowing and being justified are natural relations. Many versions of reliabilism about justification and knowledge might be seen as congenial to the naturalistic worldview? The basic idea behind these theories is that justification for a belief is a matter of being formed by a process or method that tends to get at the truth. And as long as this statistical notion of "tending to get at the truth" is viewed as naturalistic, these theories seem to construe justification as a naturalistic relation.

QU5: If one is a naturalist about epistemic justification, what will one's view be in the philosophy of language and mind about epistemic statements and the judgments they express?

On the other hand, other epistemologists insist that justification is an inextricably normative notion (perhaps intimately tied up with other paradigmatically normative notions such as believing for *reasons* or believing as one *ought* to do). Someone with this view may grant that reliability has something to do with justification but maintain that there is no specific degree of reliability that adds up to justification. This in turn leads some metaepistemologists to embrace a form of nonnaturalism with respect to justification. This is a view, just like the nonnaturalist view in metaethics, that holds that epistemic judgments purport to represent reality and some of them succeed, but the facts represented don't fit within the naturalistic worldview (see Cuneo, 2007).

Another way to motivate a kind of nonnaturalism in metaepistemology is to use Moore's open-question argument as a lens through which to view the literature responding to the **Gettier problem**. Gettier famously challenged the traditional JTB analysis of knowledge by providing clear cases of someone with a justified true belief that we wouldn't typically call knowledge. And for each tweak or complication proposed to avoid such cases, ingenious epistemologists have come up with new cases of someone meeting all of the proposed criteria but failing to have knowledge. As we discussed in Chapter 2, the most plausible way to view Moore's argument is in inductive terms: (roughly) all attempts to analyze *X* can be seen to fail to close what should be a closed question, so *X* is simple and unanalyzable. Moore applied this to the concept of goodness and considered only a few proposed analyses to confirm his semantic intuition that 'good' simply cannot mean the same as '*N*' for any proposed naturalistic *N*. In epistemology, our semantic intuitions are probably less strong. Is it an open question to ask, "She has a justified true belief that *p*, but does she *know* that *p*?" Perhaps our semantic intuitions don't settle the matter, hence the need for clear counter-examples to undermine a proposed analysis. So the post-Gettier literature could be seen as considering many proposed analyses of knowledge and showing them to fail via counter-examples meant to shore up semantic intuitions. Maybe this supports the view that knowledge is simple and unanalyzable. Indeed, some recent **knowledge-first epistemology** inspired by Williamson (2000) can be seen as diagnosing the failure of epistemologists to come up with an analysis of knowledge by the simple and unanalyzable nature of knowledge.

Metaepistemological antirealism

You'll recall, however, that Moore's open-question argument has provided as much motivation for expressivist views in metaethics as it has for nonnaturalist views. The

idea we discussed in Chapter 3 is that Moore's open questions might indicate not that ethical facts are nonnatural, but rather that proposed naturalistic analyses of ethical concepts leave out the prescriptive or endorsing element of the relevant ethical concept. A similar idea could be applied in the epistemic domain. For instance, one might think that there are various naturalistic relations between believers and propositions that are relevant to whether their beliefs count as justified or knowledge, but one might still deny any of these relations fully analyze the concepts of justification and knowledge. For one might think that those concepts have an endorsing or prescriptive element that cannot be carried by the naturalistic relations alone.

This leads us away from metaepistemological realism to an expressivist form of metaepistemological antirealism. You'll recall that *metaethical* expressivists think that ethical statements are not (solely) representations of how reality is, but rather serve primarily to express some attitude with a directive rather than descriptive direction of fit with the world. *Metaepistemological* expressivists think something similar about epistemic judgments. There are various ways to develop this thesis. One might think, with Austin (1979, 99), that claims to knowledge involve the expression of some kind of guarantee, or with Rorty (1979, 175), that claims that a belief is justified or constitutes knowledge are more comments on its status among one's peers as descriptions of the relation between mind and world.

However, the most sober and systematic forms of epistemic expressivism are developed in analogy to more contemporary forms of ethical expressivism. Indeed, one of the first places where epistemic expressivism was developed was by *ethical* expressivists. For example, Blackburn argues that, "the primary function of talking of 'knowledge' is to indicate that a judgment is beyond revision" (1998, 318). In his view, it is marked out as "beyond revision" in the sense that "no further useful investigation or thought ought to undermine the [judgment]" (1996, 87). Similarly, Gibbard argues that attributions of knowledge, such as "Joe knows there are cows on the hill," means "very roughly...that judgments like his are to be relied on" (2003, 227). This means that these judgments count as what he would call "plan-laden" rather than descriptive of reality.

So the core idea in metaepistemological expressivism is that judgments about someone's belief being justified and/or someone knowing something are not purely descriptive but in part the expression of a desire-like attitude. The reasons ethical expressivists such as Blackburn and Gibbard sought to develop this kind of view alongside their versions of expressivism about ethics are interesting for the example of how meta-level theories of different domains of normativity can seem to hang together in one big package.

Here's one reason: if you think that part of what's crucial to a concept being "normative" is that it expresses prescriptive or endorsing attitudes, and you think that evaluating actions as ethically *justified* involves the application of a normative concept, then evaluating beliefs as epistemically *justified* should also seem to involve the application of a normative concept and so the expression of an endorsing or

prescriptive attitude. Hence, the convinced expressivist about ethical normativity might quickly become a convinced expressivist about epistemic normativity, simply by recognizing the concept of *justification* as normative.

Here's another reason for ethical expressivists to be epistemic expressivists: recall that early expressivists denied the possibility of *ethical* knowledge. But this commitment flies in the face of ordinary ethical discourse, where it is a commonplace to claim to know something to be right/wrong, good/bad, etc. This is part of the reason philosophers such as Blackburn and Gibbard began to develop the "quasi-realist" program of explaining how, from within an expressivist viewpoint, we could still make sense of ordinary talk of ethical truth, belief, and knowledge. Now, if we assume, as the metaepistemological realist seems to, that knowledge is basically some kind of relation between mind and reality, but we deny that ethical facts are really "out there" in reality, then it will be mysterious how it could make sense to speak of knowing some ethical truth. If, on the other hand, one thought that talk about knowledge was itself, at least in part, a form of prescription or endorsement rather than representation, then the path is opened to an easier quasi-realist explanation of what is going on when we claim ethical knowledge.

Those are not, however, the only reasons philosophers have given for epistemic expressivism. Leaving aside metaethics for a moment, from strictly within epistemology, some philosophers have argued that the term 'knows' is somehow context-sensitive. That means, in some contexts a lot of evidence is required to claim truly to know something (e.g., as a witness in a murder trial), whereas in other contexts not much evidence is required to claim truly to know something (e.g., in recounting one's day while sharing a drink at the pub). In the past few decades, this has led to the development of various forms of epistemic **contextualism**, the view that what it takes to truly be said to know something varies with context. (This might be the context in which the putative knower finds herself or the context from which the knowledge is attributed, representing two different forms of contextualism.)

One big challenge for contextualists, however, is to make sense of intuitions of disagreement and changing one's mind when it comes to knowledge attributions made in different contexts. For example, if in the bar I say:

- (1) I know that James was at the office on Christmas day.

But then in a court of law I say:

- (2) I don't know that James was at the office on Christmas day.

Someone who heard me say both (1) and (2) might ask me why I changed my mind. But on the contextualist view, I might insist that I haven't changed my mind: when I said (1) it was true, and when I said (2) it was true too; I'm just speaking in different contexts. Similarly, you might say in a low-standards context:

- (3) Sara knows that the post office is open on Saturday.

While I might say in a high-standards context that, no, (3) is wrong:

(4) Sara doesn't know that the post office is open on Saturday.

In this kind of case, it can seem to a third-party that, in saying (3) and (4), you and I are disagreeing about whether Sara has a particular piece of knowledge. But as far as contextualism is concerned, both you and I could be speaking truly. So it is unclear where the disagreement lies.

The careful reader will have noticed that this is an analogous problem to the problem from disagreement faced by metaethical relativists. The contextualist in epistemology is a kind of relativist about knowledge attributions. She claims (roughly) that "S knows that *p*" means that S's belief that *p* is legitimate given evidential norms *N*; and she allows that *N* can vary from context to context like the ethical relativist allows that ethical norms can vary from place to place. An norm-expressivist following Gibbard (1990), however, might argue that ethical and epistemic claims are not the *representation* of what is legitimate relative to norms but rather the expression of a commitment to those norms. That is, an epistemic expressivist might reject contextualism in favor of the view that "S knows that *p*" is a vehicle for expressing commitment to some set of evidential norms. Then, the apparent variation of standards for knowledge can be explained by different evidential norms being endorsed in different contexts in which 'knows' is used. But the intuitions about changing one's mind and disagreement can be explained in terms of the sort of "disagreement in attitude" that comes along with endorsing different norms. Anyway, that's the sort of argument for epistemic expressivism pursued in more detail in Chrisman (2007, 2012). It doesn't depend on a commitment to ethical expressivism.

We've quickly surveyed some arguments for expressivism about epistemic normativity as a form of metaepistemological antirealism. It's worth asking whether the other main antirealist tradition in metaethics—error theory/fictionalism—offers an alternative form of metaepistemological antirealism. Traditionally, there has been a kind of **radical skepticism** which might be interpreted as a form of metaepistemological error theory. It is not enough to argue that we don't know much of what we think we know, for as long as one holds that there is knowledge, one will think some possible knowledge attributions are true. But if one thought that the very idea of knowledge (and/or justified belief) contained some kind of error, such that there couldn't be knowledge (and/or justified belief), then one would be a metaepistemological error-theorist. Some classical stoical arguments for skepticism have this character. And more recently, error-theorists about normativity in general have argued against the reality of reasons, both practical and epistemic (see Olson, 2011). This too would generate a form of metaepistemological error theory.

We have by no means exhausted the possibilities for metaepistemological theorizing, or even gotten a good lie of the land in this brief survey. But I hope the discussion has helped you to see some of the issues that are parts of the debate actively underway in the contemporary literature.

• CONCLUSION

This final chapter has opened up the discussion turning from metaethics to other cognate areas.

First we considered normativity in general, conceiving of ethics as a species of this larger genus. This relatively new area of philosophical inquiry is called metanormative theory. As a partial review of the theoretical traditions considered earlier in the book and as part of surveying some of the initial theoretical possibilities in metanormative theory, we considered what each of nonnaturalism, expressivism, error theory/fictionalism, and naturalism would look like when transposed to all-things-considered normativity. What we found is that many but not all of the arguments for and against these views would apply to the more general normativity.

Then we drilled down to the epistemic domain, as this is one of the species of normativity that has received the most attention in recent literature. This provides a good example of how the sorts of meta-level metaphysical, epistemological, linguistic, and psychological issues that frame metaethical theorizing can arise in other normative domains.

• CHAPTER SUMMARY

- Metanormative theory concerns questions about the status of *normative* judgments, much like metaethics concerns questions about the status of *ethical* judgments.
- It is possible to develop forms of nonnaturalism, expressivism, error theory/fictionalism, and naturalism about all-things-considered normativity.
- Proponents of these views about ethics will often think they are also the correct views to take about normativity in general, as many of the positive arguments in favor of those theories carry over to metanormative theory.
- However, many of the counter-arguments also carry over.
- Metaepistemology is the analog of metaethics, where ethics and epistemology are conceived as two separate normative domains.
- Much traditional epistemic theorizing in the attempt to analyze knowledge can be interpreted as a form of metaepistemological realism. However, it is unclear whether it should be interpreted as a form of nonnaturalism or naturalism.
- There are forms of metaepistemological antirealism, especially in the form of expressivist views about justification and knowledge.

• STUDY QUESTIONS

- 1 What are two arguments for a metaethical view that don't seem as plausible when transposed into arguments for a metanormative view? Explain your answer.
- 2 What is an objection to a metaethical view that loses some of its force if applied to the analogous metanormative view?

- 3 Is the “existential viewpoint” mentioned a kind of error theory about all-things-considered normativity?
- 4 In what sense is *justification* a normative concept? Is this the same sense of “normative” as involved in the assumption that ethical right/wrong are normative concepts?
- 5 What are some of the factual conditions that must apply for someone to know something? Why might someone think this doesn’t exhaust the concept of knowledge?
- 6 What are some of the potential problems with metaepistemological expressivism?

• FURTHER RESOURCES

- Chrisman, Matthew. 2012. “Epistemic Expressivism,” *Philosophy Compass* 7 (2): 118–26. [Survey of many of the arguments for and against epistemic expressivism.]
- Cuneo, Terence. 2007. *The Normative Web*. Oxford University Press. [Research monograph arguing that epistemic and ethical normativity hang together, deserving a nonnaturalist treatment.]

• ANSWERS TO QUESTIONS OF UNDERSTANDING

- QU1: Some general normative terms are *reason*, *should*, *ought*, *rational*, *right/wrong*, *permissible/impermissible*. Sometimes we also include *better/worse* and *virtue/vice*.
- QU2: Natural and supernatural facts.
- QU3: If there are two situations identical in all nonnormative respects, then the supervenience intuition is that what one ought all-things-considered to do in the first situation must be the same as what one ought all-things-considered to do in the second situation.
- QU4: They would say that this claim too is not an assertion but rather some other speech-act, and as such it can be perfectly in order even if the sentence used to make it is literally false.
- QU5: Naturalists are typically representationalist success-theorists in the philosophy of language, which means that they hold that epistemic statements purport and succeed in representing reality. This goes hand-in-hand with cognitivist views about the type of mental states expressed by epistemic statements. That is, naturalists will usually think that such statements express beliefs.

• WORKS CITED

- Austin, J. L. 1979. *Philosophical Papers*. Oxford: Oxford University Press.
- Blackburn, Simon. 1996. “Securing the Nots: Moral Epistemology for the Quasi-Realist.” In *Moral Knowledge: New Readings in Moral Epistemology*, edited by W. and M. Timmons Sinnott-Armstrong. Oxford; New York: Oxford University Press.

- . 1998. *Ruling Passions: A Theory of Practical Reasoning*. New York: Oxford University Press.
- Chrisman, Matthew. 2007. "From Epistemic Contextualism to Epistemic Expressivism." *Philosophical Studies* 135: 225–54.
- Chrisman, Matthew. 2012. "Epistemic Expressivism," *Philosophy Compass* 7 (2): 118–26.
- Cuneo, Terence. 2007. *The Normative Web*. New York: Oxford University Press.
- Gibbard, Allan. 1990. *Wise Choices, Apt Feelings: A Theory of Normative Judgment*. Cambridge, MA: Harvard University Press.
- . 2003. *Thinking How to Live*. Cambridge, MA: Harvard University Press.
- Olson, Jonas. 2010. "In Defence of Moral Error Theory." In *New Waves in Metaethics*, edited by M. Brady. New York: Palgrave Macmillan.
- . 2011. "Error Theory and Reasons for Belief." In *Reasons for Belief*, edited by A. Steglich-Petersen and A. Reisner, 75–93. Cambridge: Cambridge University Press.
- Rorty, Richard. 1979. *Philosophy and the Mirror of Nature*. Princeton, NJ: Princeton University Press.
- Williamson, Timothy. 2000. *Knowledge and its Limits*. Oxford; New York: Oxford University Press.